

CSCI3360 Introduction to Computational Complexity (2026 Fall)

Xiao Liang

 <https://xiao-liang.github.io>

Department of Computer Science and Engineering
The Chinese University of Hong Kong

(Updated on September 9, 2026)

Disclaimer: These lecture notes originate from the 2026-Fall offering of the course. They have been lightly edited for accuracy and completeness, and may still contain typographical or other errors.

Contents

1	The of Computation: Turing Machines	1
1.1	Before Turing Machines: Restricted Models of Computation	1
1.1.1	Finite Automata	1
1.1.2	Pushdown Automata	2
1.2	Turing Machines	4
1.2.1	Informal Description	4
1.2.2	Formal Definition	4
1.2.3	A Turing Machine for Palindromes	7
1.2.4	Running Example: $x = \text{radar}$	8
1.2.5	Going Beyond Pushdown Automata	10
1.3	Robustness of the Turing-Machine Model	11
1.3.1	Three Tapes Versus One Tape	11
1.3.2	One-Way Versus Bidirectional Tapes	13
1.4	Universal Turing Machines	15
1.4.1	Machines as Data	15
1.4.2	Definition and Basic Simulation	16
1.4.3	Efficient Universal Simulation	17
1.4.4	*Proof of Theorem 1.4.2	18
	Bibliography	19

Chapter 1

The of Computation: Turing Machines

Computational complexity asks not only whether a problem can be solved, but also how many computational resources are required to solve it. Before we can measure resources such as time and memory, however, we must specify what counts as a computation. In other words, we need a mathematical model of an algorithm.

We begin by recalling two models from formal-language theory: finite automata and push-down automata, which were already covered in **CSCI3130 Formal Languages and Automata Theory**. These models are useful precisely because their memory is restricted, but the same restrictions prevent them from expressing general computation. We then introduce the Turing machine, work through algorithms for recognizing palindromes and a non-context-free language, and examine the robustness of the model. Finally, we introduce universal Turing machines and discuss the Church–Turing thesis and its complexity-theoretic extension.

To study the complexity of computation, we must first define a model of computation. The model must be precise enough for mathematical proofs and expressive enough to represent general algorithms.

1.1 Before Turing Machines: Restricted Models of Computation

Nothing from [Section 1.1](#) will be tested on quizzes or exams!

1.1.1 Finite Automata

Recall that a deterministic finite automaton has a finite set of states and reads its input from left to right. At every step, its next state is determined by its current state and the next input symbol.

Definition 1.1.1 (Deterministic Finite Automaton). A deterministic finite automaton, or *DFA*, is a tuple

$$A = (Q, \Sigma, \delta, q_0, F),$$

where:

- Q is a finite set of states;
- Σ is a finite input alphabet;
- $\delta : Q \times \Sigma \rightarrow Q$ is the transition function;
- $q_0 \in Q$ is the start state; and
- $F \subseteq Q$ is the set of accepting states.

The automaton reads its input only once, from left to right. It accepts an input $x \in \Sigma^*$ if the state reached after reading all of x belongs to F . ◇

The current state of a DFA contains all the information that the machine retains about the portion of the input that it has already read. Since Q is finite, a DFA has only a constant amount of memory, independent of the input length.

Example 1.1.1 (A Language Recognized by a DFA). *Consider the language*

$$L_{\text{even}} := \{x \in \{0,1\}^* : x \text{ contains an even number of 1's}\}.$$

A two-state DFA recognizes this language. One state records that the number of 1's seen so far is even, and the other records that it is odd. Reading a 1 switches states, while reading a 0 leaves the state unchanged. This is illustrated in Figure 1.1.




Figure 1.1: A DFA recognizing binary strings containing an even number of 1's.

The boundary of finite-state computation

Finite automata recognize exactly the *regular languages*. They are well suited to problems in which only a bounded amount of historical information is needed.

For example, finite automata are used in the lexical-analysis stage of a compiler. A lexer may recognize identifiers, integer literals, punctuation, and keywords using regular expressions and finite-state machinery.

The limitation is that a DFA cannot store an unbounded counter. Consequently, it cannot recognize the language

$$L_{\text{eq}} := \{0^n 1^n : n \geq 0\}. \quad (1.1)$$

To recognize L_{eq} , a machine must compare an arbitrarily large number of 0's with an arbitrarily large number of 1's.

A finite set of states cannot remember the required number for all possible input lengths.

1.1.2 Pushdown Automata

A pushdown automaton extends a finite automaton by providing it with a stack. The stack can grow with the input, so a pushdown automaton can remember an unbounded amount of information. However, that information can be accessed only in last-in, first-out order.

Definition 1.1.2 (Deterministic Pushdown Automaton). *A deterministic pushdown automaton (DPDA) is a 6-tuple*

$$M = (Q, \Sigma, \Gamma, \delta, q_0, F),$$

where:

1. Q is a finite set of states;
2. Σ is the finite input alphabet;

3. Γ is the finite stack alphabet;
4. $\delta: Q \times \Sigma_\varepsilon \times \Gamma_\varepsilon \rightarrow Q \times \Gamma_\varepsilon$ is the transition function; Here,

$$\Sigma_\varepsilon = \Sigma \cup \{\varepsilon\} \quad \text{and} \quad \Gamma_\varepsilon = \Gamma \cup \{\varepsilon\},$$

where ε denotes the empty string.

5. $q_0 \in Q$ is the start state; and
6. $F \subseteq Q$ is the set of accept states.

The automaton reads its input only once, from left to right. It accepts an input $x \in \Sigma^*$ if the state reached after reading all of x belongs to F .

◇

Remark 1.1.1 (On Nondeterminism). There is a more powerful version called a nondeterministic pushdown automaton (NPDA), in which the transition function is allowed to return a set of possible next states and stack symbols. Deterministic and nondeterministic pushdown automata do not have the same expressive power. Every language recognized by a DPDA can be recognized by an NPDA, but the converse does not hold. More precisely,

$$\text{DCFL} \subsetneq \text{CFL}.$$

Thus, NPDAs recognize all context-free languages, whereas DPDAs recognize only the proper subclass of deterministic context-free languages.

Think of the stack as restricted working memory. The machine can access the most recently stored symbol, but it cannot directly inspect a symbol buried in the middle of the stack.

Example 1.1.2 (Recognizing $0^n 1^n$ with a Stack). A DPDA can recognize L_{eq} from [Equation \(1.1\)](#) as follows:

1. For each 0 in the first part of the input, push one marker onto the stack.
2. When the first 1 is encountered, begin popping one marker for each 1.
3. Accept if the input ends exactly when the stack returns to its initial marker.

For the input 000111, the stack first grows by three markers and then shrinks by three markers. For 00111, the machine attempts to pop a third marker after only two have been stored, so it rejects.

Nondeterministic pushdown automata recognize exactly the *context-free languages*. In compiler design, they provide the mathematical basis for parsing nested syntactic structures, such as balanced parentheses, arithmetic expressions, and recursively nested blocks.

The boundary of stack-based computation

Although a stack is more powerful than finite-state memory, one stack does not provide general-purpose random-access memory. A standard example beyond the power of pushdown automata is

$$L_{\text{three}} := \{0^n 1^n 2^n : n \geq 0\}. \quad (1.2)$$

A machine recognizing this language must verify that three independently delimited blocks have the same length.

A stack can conveniently match the 0's against the 1's by pushing for the first block and popping for the second. After those comparisons, however, the stored count has been destroyed, leaving no copy with which to check the 2's.

Fact 1.1.1 (Pushdown-Automaton Boundary). *The language L_{three} in Equation (1.2) is not context-free. Therefore, no pushdown automaton recognizes it.*

The statement in Fact 1.1.1 can be proved using the pumping lemma for context-free languages or Ogden's lemma. The important point for us is conceptual: the stack provides unbounded but highly structured memory, and that structure imposes a genuine limit on the languages a PDA can recognize.

Takeaway

Finite automata and pushdown automata remain important because their restrictions make them suitable for particular applications and allow precise classifications of languages. They are not, however, expressive enough to represent every algorithm that we would normally regard as executable by a computer. This motivates the Turing-machine model.

1.2 Turing Machines

Our next goal is to formalize a general-purpose model of computation. To this end, we introduce *Turing machines*, which may be viewed roughly as finite automata augmented with unbounded memory and the ability to move back and forth over their input, rather than being restricted to a single left-to-right pass.

1.2.1 Informal Description

An informal algorithm consists of a fixed collection of mechanical rules that can be applied repeatedly. A typical step may perform the following operations:

- read symbols from the input or working memory;
- write symbols to working memory;
- update a finite control state;
- move to nearby memory locations; and
- halt with an output.

A Turing machine turns this description into a precise mathematical object. Each tape is divided into cells, and each cell stores one symbol from a finite alphabet. A tape head reads the symbol in the current cell and may move one position left or right.

The tapes in our definition are unbounded to the right and have a leftmost cell. This is a mathematical abstraction: an execution lasting T steps can visit at most $T + 1$ cells on any tape, so every finite computation uses only a finite portion of each tape.

1.2.2 Formal Definition

Definition 1.2.1 (Three-Tape Turing Machine). *A three-tape Turing machine is a tuple*

$$M = (\Gamma, Q, \delta),$$

with the following components:

- Γ is a finite set of symbols that M 's tape can contain. We call it the alphabet.
We assume that it contains a blank symbol $\sqcup$, a left-end marker $\triangleright$, and a set of some characters¹ (e.g., the English alphabet, or the binary bits $\{0, 1\}$).
- Q is a finite set of states containing a start state q_{start} and a halting state q_{halt} .
- The transition function is

$$\delta : Q \times \Gamma^3 \longrightarrow Q \times \Gamma^2 \times \{L, S, R\}^3. \quad (1.3)$$

The first tape is a read-only input tape. The second tape is a read/write work tape, and the third tape is a read/write output tape. $\diamond$

Let us unpack the transition function in Equation (1.3). Suppose

$$\delta(q, a, b, c) = (q', b', c', d_1, d_2, d_3).$$

This transition instructs the machine to:

1. change its state from q to q' ;
2. leave the input symbol a unchanged; (recall that the input tape is read-only;)
3. replace the work-tape symbol b by b' ;
4. replace the output-tape symbol c by c' ; and
5. move the three heads according to

$$d_1, d_2, d_3 \in \{L, S, R\},$$

where L , S , and R mean move left, stay in place, and move right, respectively.

If a head is at the left-end marker and is instructed to move left, it remains in place.

Definition 1.2.2 (Start Configuration). *On input $x = x_1x_2 \cdots x_n \in \Gamma^*$, the start configuration of a three-tape Turing machine is defined as follows:*

- the input tape contains

$$\triangleright x_1x_2 \cdots x_n \sqcup \sqcup \cdots ;$$

- the work and output tapes each contain

$$\triangleright \sqcup \sqcup \sqcup \cdots ;$$

- all three heads are placed at their left-end markers; and
- the finite control is in state q_{start} .

$\diamond$

The machine proceeds by repeatedly applying its transition function. If it enters q_{halt} , its computation stops. The finite string written immediately to the right of the output tape's left-end marker is interpreted as the machine's output.

¹Note that the choice of characters is arbitrary. But it is without loss of generality that we only consider binary alphabets, because any finite alphabet can be encoded as a binary string. One potential issue is the size of encoding may grow. But as we will see, this does not affect the computational power of Turing machines, at least to the extent that this course is concerned.

Definition 1.2.3 (Computing a Function and Running Time). *Let $f : \{0, 1\}^* \rightarrow \{0, 1\}^*$, and let $T : \mathbb{N} \rightarrow \mathbb{N}$. A Turing machine M computes f in time $T(n)$ if, for every input $x \in \{0, 1\}^*$, the machine:*

1. *starts in the start configuration on x ;*
2. *halts after at most $T(|x|)$ transitions; and*
3. *has $f(x)$ written on its output tape when it halts.*

The running time of M is the maximum number of transitions it performs over all inputs of a given length. $\diamond$

The quantity $n = |x|$ is the input length. Complexity theory studies how resource usage grows as a function of n , rather than how long a machine takes on only one particular input.

Definition 1.2.4 (Deciding a Language). *Let $L \subseteq \{0, 1\}^*$. A Turing machine M decides L if it halts on every input x and outputs*

$$M(x) = \begin{cases} 1, & x \in L, \\ 0, & x \notin L. \end{cases} \quad (1.4)$$

Equivalently, M computes the characteristic function of L . $\diamond$

Remark 1.2.1. *In complexity theory, we normally study machines that halt on every input. A machine that accepts members of a language but may run forever on nonmembers is called a recognizer rather than a decider. Unless stated otherwise, the algorithms in these notes are deciders.*

Notice that [Definition 1.2.4](#) use the concept of a language to formalize the notion of a decision problem. Do not worry if you are not familiar with the concept of a language. Let us explain it in the following example.

Example 1.2.1 (Representing a Property as a Language). *A language is a set of strings satisfying a specified property. For example, consider the property*

“The number of 1’s in the binary string is a multiple of 3.”

This property determines the language

$$L_{\text{mult-3}} := \{x \in \{0, 1\}^* : \text{the number of 1's in } x \text{ is a multiple of 3}\}.$$

If $\#_1(x)$ denotes the number of occurrences of 1 in x , then the same language can be written more concisely as

$$L_{\text{mult-3}} = \{x \in \{0, 1\}^* : \#_1(x) \equiv 0 \pmod{3}\}.$$

For instance,

$$\varepsilon, 0, 000, 111, 10101 \in L_{\text{mult-3}},$$

because these strings contain either zero or three occurrences of 1. On the other hand,

$$1, 10, 110, 1111 \notin L_{\text{mult-3}}.$$

Notice that 0 is a multiple of 3, so every string containing no 1’s, including the empty string ε , belongs to the language.

This language corresponds to the following decision problem:

Input: A binary string x .

Question: Is the number of 1's in x a multiple of 3?

The answer is “yes” exactly when $x \in L_{\text{mult-3}}$, and it is “no” exactly when $x \notin L_{\text{mult-3}}$.

Remark 1.2.2 (Decision Problems as Languages). *Every decision problem whose instances have finite descriptions can be formalized as a language. First, encode each problem instance I as a binary string $\langle I \rangle$. The language associated with the problem is then*

$$L = \{\langle I \rangle \in \{0,1\}^* : \text{the answer for instance } I \text{ is “yes”}\}.$$

Thus, solving the decision problem is equivalent to deciding whether the encoding $\langle I \rangle$ belongs to L .

Strings in L are called yes-instances, while strings not in L are called no-instances.

Why Turing Machines?

The Turing machine provides a simple but powerful mathematical model of general computation. Its finite control represents the machine’s fixed program: the set of states and transition rules is chosen in advance and does not depend on the particular input. The work tape represents memory whose size may grow as the computation proceeds. Thus, unlike a finite automaton, a Turing machine can store an amount of information that is not bounded by a constant.

Despite having potentially unbounded memory, the machine operates locally. During one transition, it can inspect only the three currently scanned symbols, overwrite only the scanned tape cells, and move each head by at most one cell. A complicated computation must therefore be carried out as a sequence of elementary local operations. Consequently, the number of transitions gives a meaningful measure of sequential running time, while the number of tape cells used gives a meaningful measure of memory consumption.

This model is also sufficiently general to express ordinary algorithms. Data structures, intermediate results, and program configurations can all be encoded as strings on the tapes. Although such encodings may be less convenient than the data types and instructions of a programming language, they allow a Turing machine to simulate the essential steps of a computer program. In particular, a universal Turing machine can receive the description of another machine together with its input and simulate that machine, capturing the idea of a stored-program computer.

Finally, the simplicity of the model makes mathematical analysis possible. Questions about whether a problem is computable can be stated precisely in terms of whether some Turing machine decides the corresponding language. Questions about efficiency can likewise be studied by counting the machine’s transitions and tape usage. Moreover, standard variants of the model—such as machines with one tape, several tapes, or different tape alphabets—compute the same class of languages and can simulate one another with controlled overhead. The Turing machine therefore abstracts away implementation details while retaining the fundamental notions of algorithm, memory, running time, and computability.

1.2.3 A Turing Machine for Palindromes

We now use the formal model to design a familiar algorithm. The purpose of this example is not to argue that palindrome testing is difficult. Rather, it illustrates how an ordinary algorithm can be implemented using tapes, head movements, and a finite collection of states.

Definition 1.2.5 (Palindrome Language). *For an alphabet Σ , define*

$$\text{PAL}_\Sigma := \{x \in \Sigma^* : x = x^R\},$$

where x^R denotes the reversal of x .

For example,

$$\text{radar}^R = \text{radar}, \quad \text{complexity}^R = \text{ytixelpmoc}.$$

Thus, $\text{radar} \in \text{PAL}_\Sigma$, while $\text{complexity} \notin \text{PAL}_\Sigma$. ◇

Algorithm 1.2.1: Three-Tape Palindrome Algorithm

On input $x = x_1x_2 \cdots x_n$:

1. Scan the input from left to right, copying each symbol to the work tape.
2. When the input head reaches the first blank after x , move the work head left so that it scans x_n .
3. Return the input head to the beginning of x , while leaving the work head in place.
4. Repeatedly compare the input symbols from left to right with the work-tape symbols from right to left.
5. If any pair differs, write 0 on the output tape and halt.
6. If all pairs agree, write 1 on the output tape and halt.

1.2.4 Running Example: $x = \text{radar}$

Let us follow the machine through its major phases.

Phase 1: Copy the input. Initially, the relevant portions of the tapes are

```

input   : ▷ r a d a r □,
work    : ▷ □ □ □ □ □ □,
output  : ▷ □.
```

where the **red color** indicates the current position of each head.

After the copying phase, the work tape contains a copy of the input:

```

input   : ▷ r a d a r □,
work    : ▷ r a d a r □,
output  : ▷ □.
```

At this point, both heads have reached the blank immediately following the word.

Phase 2: Position the two heads. The work head moves left once and therefore scans the last symbol, **r**. The input head moves left until it reaches ▷, and then moves right once to the first symbol. The heads now scan opposite ends of the word:

```

input   : ▷ r a d a r □,
work    : ▷ r a d a r □,
output  : ▷ □.
```

Phase 3: Compare in opposite directions. The sequence of comparisons is shown in

Table 1.1.

Comparison	Input-tape symbol	Work-tape symbol	Result
1	r	r	Match
2	a	a	Match
3	d	d	Match
4	a	a	Match
5	r	r	Match

Table 1.1: The comparison phase of the palindrome machine on input **radar**. The input head moves right while the work head moves left.

After the fifth successful comparison, the input head reaches the blank following the input and the work head reaches the left-end marker. The machine writes 1 to the output tape and halts.

Theorem 1.2.1 (Correctness and Running Time of Palindrome Testing). *The machine in Algorithm 1.2.1 decides PAL_Σ in time $3n + O(1)$, where n is the input length.* $\diamond$

Proof. Let the input be $x = x_1x_2 \cdots x_n$.

The copying phase writes x_i into the i -th work-tape cell for every $i \in \{1, \dots, n\}$. Therefore, after copying, the work tape contains exactly $x_1x_2 \cdots x_n$.

The machine then positions the input head over x_1 and the work head over x_n . After $i - 1$ successful comparisons, the input head scans x_i , while the work head scans x_{n-i+1} . Thus, the i -th comparison checks whether

$$x_i = x_{n-i+1}.$$

The machine outputs 1 exactly when this equality holds for every i . This condition is equivalent to $x = x^R$, so the algorithm is correct.

The machine makes $3n + O(1)$ transitions in total:

- $n + O(1)$ transitions to copy the input;
- $n + O(1)$ transitions to return the input head to the beginning; and
- $n + O(1)$ transitions to compare the two copies.

In each phase, the n term accounts for the transitions that process or traverse the n input symbols. The $O(1)$ term accounts for a fixed number of additional transitions, such as detecting an endmarker, changing states, repositioning a head by a fixed number of cells, and entering an accepting or rejecting state. It is not hard to see that the number of these additional transitions is bounded by a constant independent of n . ■

Remark 1.2.3. *We remark that palindromes can actually be recognized by a nondeterministic pushdown automaton. Therefore, palindrome recognition is not being used here to separate Turing machines from all pushdown automata. It is used as a simple example of Turing-machine design. The separation from pushdown automata will instead use L_{three} .*

Takeaway

The work tape allows the machine to store a copy of its input and revisit that information in a different order. This is an elementary example of a general pattern: a Turing machine can reorganize and repeatedly access information rather than being restricted to one left-to-right pass or one stack discipline.

1.2.5 Going Beyond Pushdown Automata

We now return to the language

$$L_{\text{three}} = \{0^n 1^n 2^n : n \geq 0\},$$

which no pushdown automaton recognizes. We show that a Turing machine can decide this language.

Algorithm 1.2.2: Deciding L_{three}

On input x :

1. Verify that x has the form $0^*1^*2^*$. Reject if the symbols appear in any other order.
2. While scanning the block of 0's, write one marker symbol “X” on the work tape for each 0.
3. At the beginning of the block of 1's, return the work head to its first marker.
4. For each input symbol 1, move the work head right by one marker. Reject if the input has more 1's than there are markers.
5. At the end of the 1-block, reject unless the work head has passed exactly all the markers.
6. Return the work head to its first marker.
7. For each input symbol 2, move the work head right by one marker. Reject if the input has more 2's than there are markers.
8. At the end of the input, accept if and only if the work head has passed exactly all the markers.

Although [Algorithm 1.2.2](#) is stated at a high level, each instruction can be expanded into ordinary Turing-machine transitions. We illustrate this for two representative steps. Let $\triangleright$ denote a left-endmarker and let $\sqcup$ denote the blank symbol.

- First, the machine can verify that the input has the form $0^*1^*2^*$ using only its finite control. It uses three states, corresponding to the three possible blocks. In the first state it reads 0's; upon seeing the first 1, it enters the second state; and upon seeing the first 2, it enters the third state. A 0 seen in either the second or third state, or a 1 seen in the third state, causes immediate rejection. Reaching the end of the input without such a violation confirms that the symbols occur in the required order.
- Second, consider Steps 3–5. After Step 2, the relevant part of the work tape has the form

$$\triangleright X^n \sqcup \sqcup \dots$$

To return to the first marker, the work head repeatedly moves left until it sees $\triangleright$, and then moves one cell to the right. For each input symbol 1, the machine checks that the work head is scanning an X and then moves the work head one cell to the right. If it sees $\sqcup$ while another 1 remains, then there are more 1's than 0's, so it rejects. When the input head reaches the first 2, the machine checks the work-tape symbol currently under the work head. If that symbol is X , then some markers remain and there are fewer 1's than 0's; if it is $\sqcup$, then exactly all n markers have been traversed. Thus, these elementary tests implement the high-level instruction that the numbers of 0's and 1's must be equal.

The comparison with the 2's is implemented in exactly the same way: the work head first returns to the left-endmarker and then traverses one X -cell for each 2.

These examples illustrate a general principle: phrases such as “verify the order,” “return to the first marker,” and “check that all markers have been passed” abbreviate finite sequences of transitions that inspect one symbol, change state, and move a tape head by one cell at a time.

Example 1.2.2 (Input 000111222). *The machine first records three markers:*

$$\triangleright XXX \square \square \dots$$

It then uses these markers to count the three 1's. After returning to the beginning of the markers, it uses the same three markers again to count the three 2's. Both comparisons succeed, so the machine accepts.

For the input 0011122, the machine encounters a third 1 after it has already passed its two stored markers. It therefore rejects.

Proposition 1.2.1. *Algorithm 1.2.2 is an $O(n)$ -time (i.e., linear-time) Turing machine that decides L_{three} .*

Proof. The algorithm in Algorithm 1.2.2 accepts only strings of the form $0^a 1^b 2^c$. It writes exactly a markers and accepts only if both the 1-block and the 2-block contain exactly one symbol per marker. Therefore, it accepts exactly when

$$a = b = c.$$

This is precisely the condition for membership in L_{three} .

The algorithm is in linear time: each input symbol is scanned a constant number of times, and the work-tape head makes only a constant number of complete traversals of the n stored markers. ■

Turing machines are strictly more expressive than finite automata and pushdown automata: they can decide languages, such as $\{0^n 1^n 2^n : n \geq 0\}$, that neither of the earlier models can recognize.

1.3 Robustness of the Turing-Machine Model

Many details of Definition 1.2.1 may appear arbitrary. Why three tapes rather than one? Why are tapes unbounded only to the right? Why is the input tape read-only? We now show that several natural variations produce the same class of computable functions. The simulations may change running time, but they do not change what can ultimately be computed.

1.3.1 Three Tapes Versus One Tape

Definition 1.3.1 (Single-Tape Turing Machine). *A single-tape Turing machine has one read/write tape that initially contains*

$$\triangleright x \square \square \dots$$

The same tape is used for the input, working memory, and output. ◇

Lemma 1.3.1 (A Three-Tape Machine Simulates a Single-Tape Machine). *If a single-tape Turing machine M computes a function in time $T(n)$, then a three-tape Turing machine computes the same function in time $O(T(n))$.* $\diamond$

Proof. The three-tape machine simply uses the work tape exactly as M uses its single tape, storing M 's simulated state in its own finite control.

Each transition of M is simulated by a constant number of transitions. When M halts, the simulator copies the designated output to its output tape. Thus, apart from the initial and final linear-time copies, the simulation has constant overhead per transition. $\blacksquare$

The reverse direction requires more work because one physical tape must represent three independently moving tapes.

Lemma 1.3.2 (A Single-Tape Machine Simulates a Three-Tape Machine). *Suppose a three-tape Turing machine M runs in time $T(n)$. There is a single-tape Turing machine $\widetilde{M}$ that computes the same function in time*

$$O(T(n)^2).$$

$\diamond$

Proof. Let Γ be the tape alphabet of M , including the blank symbol $\sqcup$, and write $g = |\Gamma|$. We assume that tape 1 initially contains the input, tapes 2 and 3 are initially blank, and all three heads begin at their designated starting cells.

Encoding the three tapes. The machine $\widetilde{M}$ stores the three simulated tapes using its single physical tape. An encoded physical cell has the form

$$((a_1, h_1), (a_2, h_2), (a_3, h_3)),$$

where $a_i \in \Gamma$ is the symbol stored in that cell of simulated tape i , and

$$h_i \in \{0, 1\}$$

indicates whether the head of tape i is currently scanning that cell.

Each component (a_i, h_i) within the cell $((a_1, h_1), (a_2, h_2), (a_3, h_3))$ is called a *track*. Conceptually, the three tracks are three aligned sequences laid on top of one another, with the i -th track recording the contents and head position of the i -th simulated tape.

Thus, for example, the encoded cell

$$((a_1, 0), (a_2, 1), (a_3, 0))$$

stores the symbols $(a_1, 0)$, $(a_2, 1)$, $(a_3, 0)$ on the first, second, and third tracks, respectively. The head bits indicate that the second simulated head is located at this position, whereas the first and third simulated heads are not.

For each track there are $2g$ possible symbol-marker pairs:

$$|\Gamma \times \{0, 1\}| = 2g.$$

Consequently, the number of possible three-track (i.e., one-cell) symbols is

$$(2g)^3 = 8g^3.$$

The tape alphabet may also retain a separate copy of the original input alphabet for input cells that have not yet been encoded, adding at most g symbols. In addition, the simulator may use a constant number of special symbols for bookkeeping. Hence $\widetilde{M}$ can be given a tape alphabet

of size at most

$$8g^3 + g + O(1).$$

In particular, the simulated alphabet remains finite. Notice that the head markers must be attached independently to the three components; merely adding one marked copy of each tuple would not suffice, since several simulated heads may occupy the same physical position.

At every point in the simulation, exactly one cell has $h_i = 1$ for each $i \in \{1, 2, 3\}$. The finite control of $\widetilde{M}$ also stores the current state of M .

To simulate one transition of M , the machine $\widetilde{M}$ performs two sweeps:

1. It sweeps from left to right, finds the three marked symbols, and stores their values in its finite control. Once it has the current state and the three scanned symbols, it can determine the next state, the symbols to write, and the directions to move each head; this is done by invoking M 's transition function.
2. It then sweeps back from right to left, updating the three symbols and moving the three head markers as required.

The main take way is: to simulate one transition of M , the machine $\widetilde{M}$ must sweep across the entire encoded region of its tape twice. Since M is a $T(n)$ -time machine, the encoded region has length $O(T(n))$. Therefore, simulating one transition of M takes $O(T(n))$ time. Simulating all $T(n)$ transitions consequently takes

$$T(n) \cdot O(T(n)) = O(T(n)^2).$$

■

Combining [Lemma 1.3.1](#) and [Lemma 1.3.2](#) shows that the two models compute the same functions, up to a polynomial slowdown:

Theorem 1.3.1 (Tape-Number Robustness). *A function is computable by a three-tape Turing machine if and only if it is computable by a single-tape Turing machine. The efficiency is up to a quadratic relationship in running time.* ◇

Remark 1.3.1. *The quadratic overhead in [Lemma 1.3.2](#) is an upper bound for a straightforward simulation, not a claim that every multi-tape computation inherently requires a quadratic slowdown. More sophisticated simulations can achieve smaller overhead.*

1.3.2 One-Way Versus Bidirectional Tapes

Definition 1.3.2 (Bidirectional Turing Machine). *A bidirectional Turing machine is a Turing machine whose tape cells are indexed by all integers:*

$$\dots, -3, -2, -1, 0, 1, 2, 3, \dots$$

Thus, its tape is unbounded in both directions. ◇

A standard one-way tape has cells indexed by $0, 1, 2, \dots$. To simulate a bidirectional tape, we fold it around position 0. See [Figure 1.2](#) for an illustration.

Lemma 1.3.3 (Bidirectional Tapes Can Be Folded). *Let M be a Turing machine whose tapes are infinite in both directions. If M computes a function f in time $T(n)$, then there is a standard Turing machine $\widetilde{M}$, whose tapes are infinite only to the right, that computes f in time $O(T(n))$.* ◇




Figure 1.2: Folding a bidirectional tape between positions -1 and 0 to simulate it using a one-way-infinite tape. Adapted from [AB09].

Proof. We describe how to replace each bidirectional tape of M with a standard tape that has a left endpoint. Because the number of tapes is fixed, it suffices to explain the construction for one tape.

Index the cells of M 's bidirectional tape by the integers:

$$\dots, -3, -2, -1, 0, 1, 2, 3, \dots,$$

with the tape head initially at cell 0 . We “fold” this tape between cells -1 and 0 . Cell $j \geq 0$ of the corresponding standard tape stores the pair

$$(a_j, a_{-j-1}),$$

where a_i denotes the symbol in cell i of M 's tape. Thus, the first few physical cells of $\widetilde{M}$'s tape represent

Cell of $\widetilde{M}$	0	1	2	3
First component	a_0	a_1	a_2	a_3
Second component	a_{-1}	a_{-2}	a_{-3}	a_{-4} .

The first components encode the nonnegative half of M 's tape in its usual order, while the second components encode the negative half in reverse order. Again, see Figure 1.2 for an illustration.

If M has tape alphabet Γ , then each encoded cell contains an element of

$$\Gamma \times \Gamma.$$

The simulator therefore uses Γ^2 , together with a constant number of special bookkeeping symbols, as part of its tape alphabet. This is still a finite alphabet.

The finite control of $\widetilde{M}$ records, for each simulated tape, which component is currently active:

$$\mathbf{P} \quad \text{or} \quad \mathbf{N}.$$

Mode **P** means that the simulated head is on the nonnegative half of the tape and is therefore reading the first component. Mode **N** means that it is on the negative half and is reading the second component. The physical head of $\widetilde{M}$ always scans the cell containing the currently simulated tape cell.

Suppose first that the simulator is in mode **P** and its physical head is at position j . It is then simulating the head of M at position j . Its movements are translated as follows:

Move by M	Move by $\widetilde{M}$	New mode
R	R	P
$L, j > 0$	L	P
$L, j = 0$	remain at cell 0	N.

The last line represents crossing the fold: moving left from simulated cell 0 leads to simulated

cell -1 , and both cells are represented in physical cell 0.

Now suppose that the simulator is in mode N and its physical head is at position j . It is then simulating the head of M at position $-j - 1$. Since the negative half was stored in reverse order, left and right movements are interchanged:

Move by M	Move by $\widetilde{M}$	New mode
L	R	N
$R, j > 0$	L	N
$R, j = 0$	remain at cell 0	P.

The final line represents crossing the fold from simulated cell -1 to simulated cell 0.

The leftmost physical cell can be given a specially marked version of the paired alphabet, so $\widetilde{M}$ can recognize when its head is at the fold.

In summary, on the nonnegative half, head movements are simulated normally; On the negative half, left and right movements are reversed because that half is stored in reverse order. When the simulated head crosses between positions 0 and -1 , the physical head remains at the leftmost cell, and the simulator switches between the two components.

If M has tape alphabet Γ , the simulator uses paired symbols from Γ^2 . Each transition of M can therefore be simulated using only a constant number of transitions: the simulator reads and updates the appropriate component of the pair and moves its head as described above. The same construction is applied independently to each tape of M . Therefore, a $T(n)$ -step computation of M is simulated in $O(T(n))$ time. Hence, the resulting one-way-tape machine computes the same function as M . ■

What robustness does and does not say

These simulations establish that the models compute the same functions. They do not say that every choice of model gives exactly the same running time. For example, a direct three-tape palindrome algorithm runs in linear time, while some restricted single-tape implementations may require substantially more time.

Reasonable variants of the Turing-machine model agree on computability and usually agree on broad complexity distinctions such as polynomial versus exponential time. This robustness is one reason Turing machines are useful as a foundation for complexity theory.

1.4 Universal Turing Machines

1.4.1 Machines as Data

A Turing machine has a finite description: it has a finite alphabet, finitely many states, and a finite transition table.² Therefore, the description of a machine can itself be encoded as a binary string. This simple observation allows one machine to receive another machine as input and simulate it.

Let M be a Turing machine. We write

$$\langle M \rangle \in \{0, 1\}^*$$

²Note that the infinite tape is not part of the *description*; only the sentence that “a Turing machine has an infinite tape” is part of the description; this sentence has finite number of words, and thus can be encoded as a finite string.

for a binary encoding of M . The encoding contains enough information to recover:

- the tape alphabet and an encoding of each tape symbol;
- the states, including the start and halting states;
- the number and roles of the tapes; and
- every entry of the transition function.

Because all these sets are finite, such an encoding can be constructed using ordinary binary representations and delimiters.

Definition 1.4.1 (Machine Represented by a String). *For a binary string α , let M_α denote the Turing machine represented by α . Strings that are not valid encodings may be assigned to a fixed default machine, such as a machine that immediately outputs 0 and halts.* $\diamond$

This convention ensures that every binary string denotes some machine. It is also convenient to permit harmless padding, so one machine may have more than one valid encoding.

Remark 1.4.1 (Programs as Data). *The distinction between program and data is not absolute. A machine description is data when it is read by another machine, but it specifies a program when interpreted as a transition table. This principle underlies interpreters, compilers, stored-program computers, and many later arguments in complexity theory.*

1.4.2 Definition and Basic Simulation

Definition 1.4.2 (Universal Turing Machine). *A universal Turing machine is a fixed Turing machine U such that, for every machine description α and every input x ,*

$$U(\alpha, x) = M_\alpha(x), \tag{1.5}$$

whenever $M_\alpha(x)$ halts. If $M_\alpha(x)$ does not halt, then the simulation by U also need not halt. $\diamond$

The notation $U(\alpha, x)$ means that U receives an encoding of the pair (α, x) . A standard delimiter or self-delimiting encoding can be used to distinguish the two components.

The universal machine maintains encoded versions of:

- the simulated machine's transition table;
- the simulated machine's current state;
- the contents of its tapes; and
- the positions of its tape heads.

To simulate one transition, U reads the simulated state and scanned symbols, looks up the appropriate transition in the encoded table, and updates the simulated configuration.

Analogy: Virtual Machines

A useful real-world analogy is a *virtual machine*. A virtual machine may behave like a complete computer, with its own processor, memory, disks, and operating system. Nevertheless, everything needed to specify it—its configuration, virtual hardware, and software—can be stored as finite data on a physical computer.

A physical computer can read this description and simulate the virtual computer step by step. The same physical computer can simulate many different virtual machines; which machine

it simulates depends on the description it is given. Similarly, a universal Turing machine can take an encoding $\langle M \rangle$ as input and simulate the particular Turing machine M described by that encoding.

This analogy also illustrates the idea of *programs as data*. To the host computer, a virtual-machine image is simply a file containing bits. When those bits are interpreted by virtualization software, however, they describe a computer and determine its behavior. In the same way, $\langle M \rangle$ is an ordinary binary string when treated as input, but it becomes a program when interpreted as the description of a Turing machine.

1.4.3 Efficient Universal Simulation

The existence of a universal Turing machine does not immediately imply that the simulation is efficient. In this subsection, we discuss the accurate cost of simulating a machine M_α on input x using a universal machine U . We will first show that a naive simulation may introduce quadratic overhead, and then we will present a more efficient simulation that incurs only logarithmic overhead.

Theorem 1.4.1 (Quadratic Universal Turing Machine). *There exists a fixed Turing machine U such that, for every $\alpha, x \in \{0, 1\}^*$,*

$$U(\alpha, x) = M_\alpha(x)$$

whenever $M_\alpha(x)$ halts.

Furthermore, if M_α halts on x within T steps, then U halts within

$$C_\alpha T^2 \tag{1.6}$$

steps, where C_α depends on the fixed machine M_α , including its number of states, alphabet size, and number of tapes, but does not depend on x or T . $\diamond$

Before proving the theorem, let us emphasize its main strength: the input machine M_α is completely arbitrary. For example, M_α may have any fixed number of tapes and may use any finite tape alphabet. In contrast, the universal machine U is fixed once and for all; indeed, U may be taken to be a standard three-tape Turing machine with binary tape alphabet. Nevertheless, U can simulate every machine M_α with at most quadratic overhead in its running time. Thus, although the constants involved may depend on the description α , the quadratic form of the bound is uniform over all input machines.

Proof of Theorem 1.4.1. We construct the required universal Turing machine U . In fact, U will be a standard three-tape Turing machine with binary tape alphabet $\{0, 1\}$.




Figure 1.3: A fixed universal Turing machine U simulating an arbitrary machine M_α . The machine U is a standard three-tape Turing machine with binary tape alphabet, whereas M_α may have any fixed number of tapes and any finite alphabet. Adapted from [AB09, Figure 1.6].

The basic idea is straightforward. Given the description α of a Turing machine M_α and an input x , the machine U maintains an encoding of the current configuration of M_α . It then repeatedly consults α to determine the next transition and updates the encoded configuration accordingly. This process is illustrated in [Figure 1.3](#), where the “Description of M ” and “Current state of M ” segments can be placed to the beginning part of the work tape (right before the “Simulation of M ’s work tape” segment), occupying only a constant number of cells.

The main point is that this simulation remains efficient even though M_α may have multiple tapes and an arbitrary finite alphabet. This follows from two observations:

1. Let Γ be the tape alphabet of M_α . Each symbol of Γ can be represented by a binary block of length

$$\lceil \log_2 |\Gamma| \rceil.$$

Since Γ is fixed once M_α is fixed, operations on these blocks incur only constant overhead relative to the running time of M_α .

2. Suppose that M_α has k tapes. Its k tapes can be encoded on a single tape by placing their contents on separate tracks and marking the locations of their heads. As in the proof of [Lemma 1.3.2](#), one transition of M_α can be simulated by scanning the portion of the encoded tape visited so far. After t steps, this portion has length $O(t)$, so simulating the first T transitions requires

$$\sum_{t=1}^T O(t) = O(T^2)$$

time. Although k may be arbitrarily large across different machines, it is a constant for each fixed machine M_α .

The universal machine U performs these operations directly from the binary description α . Consequently, if M_α halts on x within T steps, then U simulates the computation and produces the same output in $O(T^2)$ time, with constants that may depend on α .

This finishes the proof of [Theorem 1.4.1](#). ■

By a more involved construction, one can achieve a more efficient simulation of a universal Turing machine. The following [Theorem 1.4.2](#) shows that the overhead can be reduced to $O(T \log T)$ time. However, the proof of this theorem is much more technical and out of the scope of this undergraduate-level course. We present the proof in [Section 1.4.4](#) as an optional reading for interested students.

Theorem 1.4.2 (Efficient Universal Turing Machine). *There exists a fixed Turing machine U such that, for every $\alpha, x \in \{0, 1\}^*$,*

$$U(\alpha, x) = M_\alpha(x)$$

whenever $M_\alpha(x)$ halts.

Furthermore, if M_α halts on x within $T \geq 2$ steps, then U halts within

$$C_\alpha T \log T \tag{1.7}$$

steps, where C_α depends on the fixed machine M_α , including its number of states, alphabet size, and number of tapes, but does not depend on x or T . ◇

1.4.4 *Proof of [Theorem 1.4.2](#)

Nothing from [Section 1.4.4](#) will be tested on quizzes or exams!

Bibliography

- [AB09] Sanjeev Arora and Boaz Barak. *Computational complexity: a modern approach*. Cambridge University Press, 2009. [14](#), [17](#)