

CSCI3360 Introduction to Computational Complexity (2026 Fall)

Xiao Liang

 <https://xiao-liang.github.io>

Department of Computer Science and Engineering
The Chinese University of Hong Kong

(Updated on September 11, 2026)

Disclaimer: These lecture notes originate from the 2026-Fall offering of the course. They have been lightly edited for accuracy and completeness, and may still contain typographical or other errors.

Contents

1	The of Computation: Turing Machines	1
1.1	Before Turing Machines: Restricted Models of Computation	1
1.1.1	Finite Automata	1
1.1.2	Pushdown Automata	2
1.2	Turing Machines	4
1.2.1	Informal Description	4
1.2.2	Formal Definition	5
1.2.3	A Turing Machine for Palindromes	8
1.2.4	Running Example: $x = \text{radar}$	8
1.2.5	Going Beyond Pushdown Automata	10
1.3	Robustness of the Turing-Machine Model	12
1.3.1	Three Tapes Versus One Tape	12
1.3.2	One-Way Versus Bidirectional Tapes	14
1.4	Universal Turing Machines	16
1.4.1	Machines as Data	16
1.4.2	Definition and Basic Simulation	16
1.4.3	Efficient Universal Simulation	17
1.4.4	*Proof of Theorem 1.4.2 (Part 1/2)	19
1.4.4.1	Example 1: Computing powers by repeated squaring	19
1.4.4.2	Example 2: Incrementing a binary counter	20
1.4.5	*Proof of Theorem 1.4.2 (Part 2/2)	22
1.4.5.1	The main idea: move the contents, not the simulated heads	22
1.4.5.2	Buffers and their meaning	22
1.4.5.3	Simulating one transition	23
1.4.5.4	Simulating a move to the right	23
1.4.5.5	The cost of one rearrangement	25
1.4.5.6	Why large rearrangements are rare	25
1.4.5.7	Summing the costs	26
1.5	The Church–Turing Thesis	26
1.5.1	The Standard Church–Turing Thesis	27
1.5.2	The Extended Church–Turing Thesis	27
1.5.3	Quantum Computing and the Extended Thesis	28

1.5.4	The Turing Machine as a Final Model	28
1.6	Closing Remarks	29
	Bibliography	30

Chapter 1

The of Computation: Turing Machines

Computational complexity asks not only whether a problem can be solved, but also how many computational resources are required to solve it. Before we can measure resources such as time and memory, however, we must specify what counts as a computation. In other words, we need a mathematical model of an algorithm.

We begin by recalling two models from formal-language theory: finite automata and push-down automata, which were already covered in **CSCI3130 Formal Languages and Automata Theory**. These models are useful precisely because their memory is restricted, but the same restrictions prevent them from expressing general computation. We then introduce the Turing machine, work through algorithms for recognizing palindromes and a non-context-free language, and examine the robustness of the model. Finally, we introduce universal Turing machines and discuss the Church–Turing thesis and its complexity-theoretic extension.

To study the complexity of computation, we must first define a model of computation. The model must be precise enough for mathematical proofs and expressive enough to represent general algorithms.

1.1 Before Turing Machines: Restricted Models of Computation

Nothing from [Section 1.1](#) will be tested on quizzes or exams!

1.1.1 Finite Automata

Recall that a deterministic finite automaton has a finite set of states and reads its input from left to right. At every step, its next state is determined by its current state and the next input symbol.

Definition 1.1.1 (Deterministic Finite Automaton). A deterministic finite automaton, or *DFA*, is a tuple

$$A = (Q, \Sigma, \delta, q_0, F),$$

where:

- Q is a finite set of states;
- Σ is a finite input alphabet;
- $\delta : Q \times \Sigma \rightarrow Q$ is the transition function;
- $q_0 \in Q$ is the start state; and
- $F \subseteq Q$ is the set of accepting states.

The automaton reads its input only once, from left to right. It accepts an input $x \in \Sigma^*$ if the state reached after reading all of x belongs to F . ◇

The current state of a DFA contains all the information that the machine retains about the portion of the input that it has already read. Since Q is finite, a DFA has only a constant amount of memory, independent of the input length.

Example 1.1.1 (A Language Recognized by a DFA). *Consider the language*

$$L_{\text{even}} := \{x \in \{0, 1\}^* : x \text{ contains an even number of 1's}\}.$$

A two-state DFA recognizes this language. One state records that the number of 1's seen so far is even, and the other records that it is odd. Reading a 1 switches states, while reading a 0 leaves the state unchanged. This is illustrated in Figure 1.1.




Figure 1.1: A DFA recognizing binary strings containing an even number of 1's.

The boundary of finite-state computation

Finite automata recognize exactly the *regular languages*. They are well suited to problems in which only a bounded amount of historical information is needed.

For example, finite automata are used in the lexical-analysis stage of a compiler. A lexer may recognize identifiers, integer literals, punctuation, and keywords using regular expressions and finite-state machinery.

The limitation is that a DFA cannot store an unbounded counter. Consequently, it cannot recognize the language

$$L_{\text{eq}} := \{0^n 1^n : n \geq 0\}. \quad (1.1)$$

To recognize L_{eq} , a machine must compare an arbitrarily large number of 0's with an arbitrarily large number of 1's.

A finite set of states cannot remember the required number for all possible input lengths.

1.1.2 Pushdown Automata

A pushdown automaton extends a finite automaton by providing it with a stack. The stack can grow with the input, so a pushdown automaton can remember an unbounded amount of information. However, that information can be accessed only in last-in, first-out order.

Definition 1.1.2 (Deterministic Pushdown Automaton). *A deterministic pushdown automaton (DPDA) is a 6-tuple*

$$M = (Q, \Sigma, \Gamma, \delta, q_0, F),$$

where:

1. Q is a finite set of states;
2. Σ is the finite input alphabet;

3. Γ is the finite stack alphabet;
4. $\delta: Q \times \Sigma_\varepsilon \times \Gamma_\varepsilon \rightarrow Q \times \Gamma_\varepsilon$ is the transition function; Here,

$$\Sigma_\varepsilon = \Sigma \cup \{\varepsilon\} \quad \text{and} \quad \Gamma_\varepsilon = \Gamma \cup \{\varepsilon\},$$

where ε denotes the empty string.

5. $q_0 \in Q$ is the start state; and
6. $F \subseteq Q$ is the set of accept states.

The automaton reads its input only once, from left to right. It accepts an input $x \in \Sigma^*$ if the state reached after reading all of x belongs to F .

◇

Remark 1.1.1 (On Nondeterminism). There is a more powerful version called a nondeterministic pushdown automaton (NPDA), in which the transition function is allowed to return a set of possible next states and stack symbols. Deterministic and nondeterministic pushdown automata do not have the same expressive power. Every language recognized by a DPDA can be recognized by an NPDA, but the converse does not hold. More precisely,

$$\text{DCFL} \subsetneq \text{CFL}.$$

Thus, NPDAs recognize all context-free languages, whereas DPDAs recognize only the proper subclass of deterministic context-free languages.

Think of the stack as restricted working memory. The machine can access the most recently stored symbol, but it cannot directly inspect a symbol buried in the middle of the stack.

Example 1.1.2 (Recognizing $0^n 1^n$ with a Stack). A DPDA can recognize L_{eq} from [Equation \(1.1\)](#) as follows:

1. For each 0 in the first part of the input, push one marker onto the stack.
2. When the first 1 is encountered, begin popping one marker for each 1.
3. Accept if the input ends exactly when the stack returns to its initial marker.

Such a DPDA is shown in [Figure 1.2](#). For the input 000111, the stack first grows by three markers and then shrinks by three markers. For 00111, the machine attempts to pop a third marker after only two have been stored, so it rejects.

Nondeterministic pushdown automata recognize exactly the *context-free languages*. In compiler design, they provide the mathematical basis for parsing nested syntactic structures, such as balanced parentheses, arithmetic expressions, and recursively nested blocks.

The boundary of stack-based computation

Although a stack is more powerful than finite-state memory, one stack does not provide general-purpose random-access memory. A standard example beyond the power of pushdown automata is

$$L_{\text{three}} := \{0^n 1^n 2^n : n \geq 0\}. \quad (1.2)$$

A machine recognizing this language must verify that three independently delimited blocks have the same length.




Figure 1.2: A deterministic pushdown automaton for $L = \{1^n 2^n \mid n \geq 0\}$. Each transition has the form “input, stack top $\rightarrow$ replacement.” The automaton pushes one A for every 1 and pops one A for every 2. Here, Z is the bottom-of-stack symbol, ε means that the top symbol is removed, and $\neg$ marks the end of the input. Thus, $\neg, Z \rightarrow Z$ means that the input has ended, only Z remains on the stack, and Z is left unchanged.

A stack can conveniently match the 0’s against the 1’s by pushing for the first block and popping for the second. After those comparisons, however, the stored count has been destroyed, leaving no copy with which to check the 2’s.

Fact 1.1.1 (Pushdown-Automaton Boundary). *The language L_{three} in Equation (1.2) is not context-free. Therefore, no pushdown automaton recognizes it.*

The statement in Fact 1.1.1 can be proved using the pumping lemma for context-free languages or Ogden’s lemma. The important point for us is conceptual: the stack provides unbounded but highly structured memory, and that structure imposes a genuine limit on the languages a PDA can recognize.

Takeaway

Finite automata and pushdown automata remain important because their restrictions make them suitable for particular applications and allow precise classifications of languages. They are not, however, expressive enough to represent every algorithm that we would normally regard as executable by a computer. This motivates the Turing-machine model.

1.2 Turing Machines

Our next goal is to formalize a general-purpose model of computation. To this end, we introduce *Turing machines*, which may be viewed roughly as finite automata augmented with unbounded memory and the ability to move back and forth over their input, rather than being restricted to a single left-to-right pass.

1.2.1 Informal Description

An informal algorithm consists of a fixed collection of mechanical rules that can be applied repeatedly. A typical step may perform the following operations:

- read symbols from the input or working memory;
- write symbols to working memory;
- update a finite control state;

- move to nearby memory locations; and
- halt with an output.

A Turing machine turns this description into a precise mathematical object. Each tape is divided into cells, and each cell stores one symbol from a finite alphabet. A tape head reads the symbol in the current cell and may move one position left or right.

The tapes in our definition are unbounded to the right and have a leftmost cell. This is a mathematical abstraction: an execution lasting T steps can visit at most $T + 1$ cells on any tape, so every finite computation uses only a finite portion of each tape.

1.2.2 Formal Definition

Definition 1.2.1 (Three-Tape Turing Machine). *A three-tape Turing machine is a tuple*

$$M = (\Gamma, Q, \delta),$$

with the following components:

- Γ is a finite set of symbols that M 's tape can contain. We call it the alphabet. We assume that it contains a blank symbol $\sqcup$, a left-end marker $\triangleright$, and a set of some characters¹ (e.g., the English alphabet, or the binary bits $\{0, 1\}$).
- Q is a finite set of states containing a start state q_{start} and a halting state q_{halt} .
- The transition function is

$$\delta : Q \times \Gamma^3 \longrightarrow Q \times \Gamma^2 \times \{L, S, R\}^3. \quad (1.3)$$

The first tape is a read-only input tape. The second tape is a read/write work tape, and the third tape is a read/write output tape. $\diamond$

Let us unpack the transition function in [Equation \(1.3\)](#). Suppose

$$\delta(q, a, b, c) = (q', b', c', d_1, d_2, d_3).$$

This transition instructs the machine to:

1. change its state from q to q' ;
2. leave the input symbol a unchanged; (recall that the input tape is read-only;)
3. replace the work-tape symbol b by b' ;
4. replace the output-tape symbol c by c' ; and
5. move the three heads according to

$$d_1, d_2, d_3 \in \{L, S, R\},$$

where L , S , and R mean move left, stay in place, and move right, respectively.

If a head is at the left-end marker and is instructed to move left, it remains in place.

¹Note that the choice of characters is arbitrary. But it is without loss of generality that we only consider binary alphabets, because any finite alphabet can be encoded as a binary string. On potential issue is the size of encoding may grow. But as we will see, this does not affect the computational power of Turing machines, at least to the extent that this course is concerned.

Definition 1.2.2 (Start Configuration). *On input $x = x_1x_2 \cdots x_n \in \Gamma^*$, the start configuration of a three-tape Turing machine is defined as follows:*

- the input tape contains

$$\triangleright x_1x_2 \cdots x_n \sqcup \sqcup \cdots ;$$

- the work and output tapes each contain

$$\triangleright \sqcup \sqcup \sqcup \cdots ;$$

- all three heads are placed at their left-end markers; and
- the finite control is in state q_{start} .

◇

The machine proceeds by repeatedly applying its transition function. If it enters q_{halt} , its computation stops. The finite string written immediately to the right of the output tape's left-end marker is interpreted as the machine's output.

Definition 1.2.3 (Computing a Function and Running Time). *Let $f : \{0,1\}^* \rightarrow \{0,1\}^*$, and let $T : \mathbb{N} \rightarrow \mathbb{N}$. A Turing machine M computes f in time $T(n)$ if, for every input $x \in \{0,1\}^*$, the machine:*

1. starts in the start configuration on x ;
2. halts after at most $T(|x|)$ transitions; and
3. has $f(x)$ written on its output tape when it halts.

The running time of M is the maximum number of transitions it performs over all inputs of a given length.

◇

The quantity $n = |x|$ is the input length. Complexity theory studies how resource usage grows as a function of n , rather than how long a machine takes on only one particular input.

Definition 1.2.4 (Deciding a Language). *Let $L \subseteq \{0,1\}^*$. A Turing machine M decides L if it halts on every input x and outputs*

$$M(x) = \begin{cases} 1, & x \in L, \\ 0, & x \notin L. \end{cases} \quad (1.4)$$

Equivalently, M computes the characteristic function of L .

◇

Remark 1.2.1. *In complexity theory, we normally study machines that halt on every input. A machine that accepts members of a language but may run forever on nonmembers is called a recognizer rather than a decider. Unless stated otherwise, the algorithms in these notes are deciders.*

Notice that [Definition 1.2.4](#) use the concept of a language to formalize the notion of a decision problem. Do not worry if you are not familiar with the concept of a language. Let us explain it in the following example.

Example 1.2.1 (Representing a Property as a Language). *A language is a set of strings satisfying a specified property. For example, consider the property*

“The number of 1’s in the binary string is a multiple of 3.”

This property determines the language

$$L_{\text{mult-3}} := \{x \in \{0,1\}^* : \text{the number of 1's in } x \text{ is a multiple of 3}\}.$$

If $\#_1(x)$ denotes the number of occurrences of 1 in x , then the same language can be written more concisely as

$$L_{\text{mult-3}} = \{x \in \{0,1\}^* : \#_1(x) \equiv 0 \pmod{3}\}.$$

For instance,

$$\varepsilon, 0, 000, 111, 10101 \in L_{\text{mult-3}},$$

because these strings contain either zero or three occurrences of 1. On the other hand,

$$1, 10, 110, 1111 \notin L_{\text{mult-3}}.$$

Notice that 0 is a multiple of 3, so every string containing no 1’s, including the empty string ε , belongs to the language.

This language corresponds to the following decision problem:

Input: A binary string x .

Question: Is the number of 1’s in x a multiple of 3?

The answer is “yes” exactly when $x \in L_{\text{mult-3}}$, and it is “no” exactly when $x \notin L_{\text{mult-3}}$.

Remark 1.2.2 (Decision Problems as Languages). *Every decision problem whose instances have finite descriptions can be formalized as a language. First, encode each problem instance I as a binary string $\langle I \rangle$. The language associated with the problem is then*

$$L = \{\langle I \rangle \in \{0,1\}^* : \text{the answer for instance } I \text{ is “yes”}\}.$$

Thus, solving the decision problem is equivalent to deciding whether the encoding $\langle I \rangle$ belongs to L .

Strings in L are called yes-instances, while strings not in L are called no-instances.

Why Turing Machines?

The Turing machine provides a simple but powerful mathematical model of general computation. Its finite control represents the machine’s fixed program: the set of states and transition rules is chosen in advance and does not depend on the particular input. The work tape represents memory whose size may grow as the computation proceeds. Thus, unlike a finite automaton, a Turing machine can store an amount of information that is not bounded by a constant.

Despite having potentially unbounded memory, the machine operates locally. During one transition, it can inspect only the three currently scanned symbols, overwrite only the scanned tape cells, and move each head by at most one cell. A complicated computation must therefore be carried out as a sequence of elementary local operations. Consequently, the number of transitions gives a meaningful measure of sequential running time, while the number of tape cells used gives a meaningful measure of memory consumption.

This model is also sufficiently general to express ordinary algorithms. Data structures, intermediate results, and program configurations can all be encoded as strings on the tapes. Although such encodings may be less convenient than the data types and instructions of a programming language, they allow a Turing machine to simulate the essential steps of a computer

program. In particular, a universal Turing machine can receive the description of another machine together with its input and simulate that machine, capturing the idea of a stored-program computer.

Finally, the simplicity of the model makes mathematical analysis possible. Questions about whether a problem is computable can be stated precisely in terms of whether some Turing machine decides the corresponding language. Questions about efficiency can likewise be studied by counting the machine's transitions and tape usage. Moreover, standard variants of the model—such as machines with one tape, several tapes, or different tape alphabets—compute the same class of languages and can simulate one another with controlled overhead. The Turing machine therefore abstracts away implementation details while retaining the fundamental notions of algorithm, memory, running time, and computability.

1.2.3 A Turing Machine for Palindromes

We now use the formal model to design a familiar algorithm. The purpose of this example is not to argue that palindrome testing is difficult. Rather, it illustrates how an ordinary algorithm can be implemented using tapes, head movements, and a finite collection of states.

Definition 1.2.5 (Palindrome Language). *For an alphabet Σ , define*

$$\text{PAL}_\Sigma := \{x \in \Sigma^* : x = x^R\},$$

where x^R denotes the reversal of x .

For example,

$$\text{radar}^R = \text{radar}, \quad \text{complexity}^R = \text{ytixelpmoc}.$$

Thus, $\text{radar} \in \text{PAL}_\Sigma$, while $\text{complexity} \notin \text{PAL}_\Sigma$. ◇

Algorithm 1.2.1: Three-Tape Palindrome Algorithm

On input $x = x_1x_2 \cdots x_n$:

1. Scan the input from left to right, copying each symbol to the work tape.
2. When the input head reaches the first blank after x , move the work head left so that it scans x_n .
3. Return the input head to the beginning of x , while leaving the work head in place.
4. Repeatedly compare the input symbols from left to right with the work-tape symbols from right to left.
5. If any pair differs, write 0 on the output tape and halt.
6. If all pairs agree, write 1 on the output tape and halt.

1.2.4 Running Example: $x = \text{radar}$

Let us follow the machine through its major phases.

Phase 1: Copy the input. Initially, the relevant portions of the tapes are

input	:	▷ r a d a r □,
work	:	▷ □ □ □ □ □ □,
output	:	▷ □.

where the **red color** indicates the current position of each head.

After the copying phase, the work tape contains a copy of the input:

input : ▷ **r** a d a r ⊐,
work : ▷ **r** a d a r ⊐,
output : ▷ ⊐.

At this point, both heads have reached the blank immediately following the word.

Phase 2: Position the two heads. The work head moves left once and therefore scans the last symbol, **r**. The input head moves left until it reaches ▷, and then moves right once to the first symbol. The heads now scan opposite ends of the word:

input : ▷ **r** a d a r ⊐,
work : ▷ **r** a d a **r** ⊐,
output : ▷ ⊐.

Phase 3: Compare in opposite directions. The sequence of comparisons is shown in Table 1.1.

Comparison	Input-tape symbol	Work-tape symbol	Result
1	r	r	Match
2	a	a	Match
3	d	d	Match
4	a	a	Match
5	r	r	Match

Table 1.1: The comparison phase of the palindrome machine on input **radar**. The input head moves right while the work head moves left.

After the fifth successful comparison, the input head reaches the blank following the input and the work head reaches the left-end marker. The machine writes 1 to the output tape and halts.

Theorem 1.2.1 (Correctness and Running Time of Palindrome Testing). *The machine in Algorithm 1.2.1 decides PAL_Σ in time $3n + O(1)$, where n is the input length.* ◇

Proof. Let the input be $x = x_1x_2 \cdots x_n$.

The copying phase writes x_i into the i -th work-tape cell for every $i \in \{1, \dots, n\}$. Therefore, after copying, the work tape contains exactly $x_1x_2 \cdots x_n$.

The machine then positions the input head over x_1 and the work head over x_n . After $i - 1$ successful comparisons, the input head scans x_i , while the work head scans x_{n-i+1} . Thus, the i -th comparison checks whether

$$x_i = x_{n-i+1}.$$

The machine outputs 1 exactly when this equality holds for every i . This condition is equivalent to $x = x^R$, so the algorithm is correct.

The machine makes $3n + O(1)$ transitions in total:

- $n + O(1)$ transitions to copy the input;
- $n + O(1)$ transitions to return the input head to the beginning; and
- $n + O(1)$ transitions to compare the two copies.

In each phase, the n term accounts for the transitions that process or traverse the n input symbols. The $O(1)$ term accounts for a fixed number of additional transitions, such as detecting an endmarker, changing states, repositioning a head by a fixed number of cells, and entering an accepting or rejecting state. It is not hard to see that the number of these additional transitions is bounded by a constant independent of n . ■

Remark 1.2.3. *We remark that palindromes can actually be recognized by a nondeterministic pushdown automaton. Therefore, palindrome recognition is not being used here to separate Turing machines from all pushdown automata. It is used as a simple example of Turing-machine design. The separation from pushdown automata will instead use L_{three} .*

Takeaway

The work tape allows the machine to store a copy of its input and revisit that information in a different order. This is an elementary example of a general pattern: a Turing machine can reorganize and repeatedly access information rather than being restricted to one left-to-right pass or one stack discipline.

1.2.5 Going Beyond Pushdown Automata

We now return to the language

$$L_{\text{three}} = \{0^n 1^n 2^n : n \geq 0\},$$

which no pushdown automaton recognizes. We show that a Turing machine can decide this language.

Algorithm 1.2.2: Deciding L_{three}

On input x :

1. Verify that x has the form $0^*1^*2^*$. Reject if the symbols appear in any other order.
2. While scanning the block of 0's, write one marker symbol "X" on the work tape for each 0.
3. At the beginning of the block of 1's, return the work head to its first marker.
4. For each input symbol 1, move the work head right by one marker. Reject if the input has more 1's than there are markers.
5. At the end of the 1-block, reject unless the work head has passed exactly all the markers.
6. Return the work head to its first marker.
7. For each input symbol 2, move the work head right by one marker. Reject if the input has more 2's than there are markers.
8. At the end of the input, accept if and only if the work head has passed exactly all the markers.

Although [Algorithm 1.2.2](#) is stated at a high level, each instruction can be expanded into ordinary Turing-machine transitions. We illustrate this for two representative steps. Let $\triangleright$ denote a left-endmarker and let $\sqcup$ denote the blank symbol.

- First, the machine can verify that the input has the form $0^*1^*2^*$ using only its finite control. It uses three states, corresponding to the three possible blocks. In the first state it reads 0's; upon seeing the first 1, it enters the second state; and upon seeing the first 2, it enters the third state. A 0 seen in either the second or third state, or a 1 seen in the third state, causes immediate rejection. Reaching the end of the input without such a violation confirms that the symbols occur in the required order.
- Second, consider Steps 3–5. After Step 2, the relevant part of the work tape has the form

$$\triangleright X^n \sqcup \sqcup \dots$$

To return to the first marker, the work head repeatedly moves left until it sees $\triangleright$, and then moves one cell to the right. For each input symbol 1, the machine checks that the work head is scanning an X and then moves the work head one cell to the right. If it sees $\sqcup$ while another 1 remains, then there are more 1's than 0's, so it rejects. When the input head reaches the first 2, the machine checks the work-tape symbol currently under the work head. If that symbol is X , then some markers remain and there are fewer 1's than 0's; if it is $\sqcup$, then exactly all n markers have been traversed. Thus, these elementary tests implement the high-level instruction that the numbers of 0's and 1's must be equal. The comparison with the 2's is implemented in exactly the same way: the work head first returns to the left-endmarker and then traverses one X -cell for each 2.

These examples illustrate a general principle: phrases such as “verify the order,” “return to the first marker,” and “check that all markers have been passed” abbreviate finite sequences of transitions that inspect one symbol, change state, and move a tape head by one cell at a time.

Example 1.2.2 (Input 000111222). *The machine first records three markers:*

$$\triangleright XXX \sqcup \sqcup \dots$$

It then uses these markers to count the three 1's. After returning to the beginning of the markers, it uses the same three markers again to count the three 2's. Both comparisons succeed, so the machine accepts.

For the input 0011122, the machine encounters a third 1 after it has already passed its two stored markers. It therefore rejects.

Proposition 1.2.1. *Algorithm 1.2.2 is an $O(n)$ -time (i.e., linear-time) Turing machine that decides L_{three} .*

Proof. The algorithm in Algorithm 1.2.2 accepts only strings of the form $0^a 1^b 2^c$. It writes exactly a markers and accepts only if both the 1-block and the 2-block contain exactly one symbol per marker. Therefore, it accepts exactly when

$$a = b = c.$$

This is precisely the condition for membership in L_{three} .

The algorithm is in linear time: each input symbol is scanned a constant number of times, and the work-tape head makes only a constant number of complete traversals of the n stored markers. ■

Turing machines are strictly more expressive than finite automata and pushdown automata: they can decide languages, such as $\{0^n 1^n 2^n : n \geq 0\}$, that neither of the earlier models can recognize.

1.3 Robustness of the Turing-Machine Model

Many details of Definition 1.2.1 may appear arbitrary. Why three tapes rather than one? Why are tapes unbounded only to the right? Why is the input tape read-only? We now show that several natural variations produce the same class of computable functions. The simulations may change running time, but they do not change what can ultimately be computed.

1.3.1 Three Tapes Versus One Tape

Definition 1.3.1 (Single-Tape Turing Machine). *A single-tape Turing machine has one read/write tape that initially contains*

$$\triangleright x \sqcup \sqcup \cdots .$$

The same tape is used for the input, working memory, and output. ◇

Lemma 1.3.1 (A Three-Tape Machine Simulates a Single-Tape Machine). *If a single-tape Turing machine M computes a function in time $T(n)$, then a three-tape Turing machine computes the same function in time $O(T(n))$.* ◇

Proof. The three-tape machine simply uses the work tape exactly as M uses its single tape, storing M 's simulated state in its own finite control.

Each transition of M is simulated by a constant number of transitions. When M halts, the simulator copies the designated output to its output tape. Thus, apart from the initial and final linear-time copies, the simulation has constant overhead per transition. ■

The reverse direction requires more work because one physical tape must represent three independently moving tapes.

Lemma 1.3.2 (A Single-Tape Machine Simulates a Three-Tape Machine). *Suppose a three-tape Turing machine M runs in time $T(n)$. There is a single-tape Turing machine $\widetilde{M}$ that computes the same function in time*

$$O(T(n)^2).$$

◇

Proof. Let Γ be the tape alphabet of M , including the blank symbol $\sqcup$, and write $g = |\Gamma|$. We assume that tape 1 initially contains the input, tapes 2 and 3 are initially blank, and all three heads begin at their designated starting cells.

Encoding the three tapes. The machine $\widetilde{M}$ stores the three simulated tapes using its single physical tape. An encoded physical cell has the form

$$((a_1, h_1), (a_2, h_2), (a_3, h_3)),$$

where $a_i \in \Gamma$ is the symbol stored in that cell of simulated tape i , and

$$h_i \in \{0, 1\}$$

indicates whether the head of tape i is currently scanning that cell.

Each component (a_i, h_i) within the cell $((a_1, h_1), (a_2, h_2), (a_3, h_3))$ is called a *track*. Conceptually, the three tracks are three aligned sequences laid on top of one another, with the i -th track recording the contents and head position of the i -th simulated tape.

Thus, for example, the encoded cell

$$((a_1, 0), (a_2, 1), (a_3, 0))$$

stores the symbols $(a_1, 0)$, $(a_2, 1)$, $(a_3, 0)$ on the first, second, and third tracks, respectively. The head bits indicate that the second simulated head is located at this position, whereas the first and third simulated heads are not.

For each track there are $2g$ possible symbol-marker pairs:

$$|\Gamma \times \{0, 1\}| = 2g.$$

Consequently, the number of possible three-track (i.e., one-cell) symbols is

$$(2g)^3 = 8g^3.$$

The tape alphabet may also retain a separate copy of the original input alphabet for input cells that have not yet been encoded, adding at most g symbols. In addition, the simulator may use a constant number of special symbols for bookkeeping. Hence $\widetilde{M}$ can be given a tape alphabet of size at most

$$8g^3 + g + O(1).$$

In particular, the simulated alphabet remains finite. Notice that the head markers must be attached independently to the three components; merely adding one marked copy of each tuple would not suffice, since several simulated heads may occupy the same physical position.

At every point in the simulation, exactly one cell has $h_i = 1$ for each $i \in \{1, 2, 3\}$. The finite control of $\widetilde{M}$ also stores the current state of M .

To simulate one transition of M , the machine $\widetilde{M}$ performs two sweeps:

1. It sweeps from left to right, finds the three marked symbols, and stores their values in its finite control. Once it has the current state and the three scanned symbols, it can determine the next state, the symbols to write, and the directions to move each head; this is done by invoking M 's transition function.
2. It then sweeps back from right to left, updating the three symbols and moving the three head markers as required.

The main take away is: to simulate one transition of M , the machine $\widetilde{M}$ must sweep across the entire encoded region of its tape twice. Since M is a $T(n)$ -time machine, the encoded region has length $O(T(n))$. Therefore, simulating one transition of M takes $O(T(n))$ time. Simulating all $T(n)$ transitions consequently takes

$$T(n) \cdot O(T(n)) = O(T(n)^2).$$

■

Combining [Lemma 1.3.1](#) and [Lemma 1.3.2](#) shows that the two models compute the same functions, up to a polynomial slowdown:

Theorem 1.3.1 (Tape-Number Robustness). *A function is computable by a three-tape Turing machine if and only if it is computable by a single-tape Turing machine. The efficiency is up to a quadratic relationship in running time.* ◇

Remark 1.3.1. The quadratic overhead in [Lemma 1.3.2](#) is an upper bound for a straightforward simulation, not a claim that every multi-tape computation inherently requires a quadratic slowdown. More sophisticated simulations can achieve smaller overhead.

1.3.2 One-Way Versus Bidirectional Tapes

Definition 1.3.2 (Bidirectional Turing Machine). A bidirectional Turing machine is a Turing machine whose tape cells are indexed by all integers:

$$\dots, -3, -2, -1, 0, 1, 2, 3, \dots$$

Thus, its tape is unbounded in both directions. ◇

A standard one-way tape has cells indexed by $0, 1, 2, \dots$. To simulate a bidirectional tape, we fold it around position 0. See [Figure 1.3](#) for an illustration.




Figure 1.3: Folding a bidirectional tape between positions -1 and 0 to simulate it using a one-way-infinite tape. Adapted from [\[AB09\]](#).

Lemma 1.3.3 (Bidirectional Tapes Can Be Folded). Let M be a Turing machine whose tapes are infinite in both directions. If M computes a function f in time $T(n)$, then there is a standard Turing machine $\widetilde{M}$, whose tapes are infinite only to the right, that computes f in time $O(T(n))$. ◇

Proof. We describe how to replace each bidirectional tape of M with a standard tape that has a left endpoint. Because the number of tapes is fixed, it suffices to explain the construction for one tape.

Index the cells of M 's bidirectional tape by the integers:

$$\dots, -3, -2, -1, 0, 1, 2, 3, \dots,$$

with the tape head initially at cell 0. We “fold” this tape between cells -1 and 0 . Cell $j \geq 0$ of the corresponding standard tape stores the pair

$$(a_j, a_{-j-1}),$$

where a_i denotes the symbol in cell i of M 's tape. Thus, the first few physical cells of $\widetilde{M}$'s tape represent

Cell of $\widetilde{M}$	0	1	2	3
First component	a_0	a_1	a_2	a_3
Second component	a_{-1}	a_{-2}	a_{-3}	a_{-4}

The first components encode the nonnegative half of M 's tape in its usual order, while the second components encode the negative half in reverse order. Again, see [Figure 1.3](#) for an illustration.

If M has tape alphabet Γ , then each encoded cell contains an element of

$$\Gamma \times \Gamma.$$

The simulator therefore uses Γ^2 , together with a constant number of special bookkeeping symbols, as part of its tape alphabet. This is still a finite alphabet.

The finite control of $\widetilde{M}$ records, for each simulated tape, which component is currently active:

$$\mathbf{P} \quad \text{or} \quad \mathbf{N}.$$

Mode **P** means that the simulated head is on the nonnegative half of the tape and is therefore reading the first component. Mode **N** means that it is on the negative half and is reading the second component. The physical head of $\widetilde{M}$ always scans the cell containing the currently simulated tape cell.

Suppose first that the simulator is in mode **P** and its physical head is at position j . It is then simulating the head of M at position j . Its movements are translated as follows:

Move by M	Move by $\widetilde{M}$	New mode
R	R	P
$L, j > 0$	L	P
$L, j = 0$	remain at cell 0	N.

The last line represents crossing the fold: moving left from simulated cell 0 leads to simulated cell -1 , and both cells are represented in physical cell 0.

Now suppose that the simulator is in mode **N** and its physical head is at position j . It is then simulating the head of M at position $-j - 1$. Since the negative half was stored in reverse order, left and right movements are interchanged:

Move by M	Move by $\widetilde{M}$	New mode
L	R	N
$R, j > 0$	L	N
$R, j = 0$	remain at cell 0	P.

The final line represents crossing the fold from simulated cell -1 to simulated cell 0.

The leftmost physical cell can be given a specially marked version of the paired alphabet, so $\widetilde{M}$ can recognize when its head is at the fold.

In summary, on the nonnegative half, head movements are simulated normally; On the negative half, left and right movements are reversed because that half is stored in reverse order. When the simulated head crosses between positions 0 and -1 , the physical head remains at the leftmost cell, and the simulator switches between the two components.

If M has tape alphabet Γ , the simulator uses paired symbols from Γ^2 . Each transition of M can therefore be simulated using only a constant number of transitions: the simulator reads and updates the appropriate component of the pair and moves its head as described above. The same construction is applied independently to each tape of M . Therefore, a $T(n)$ -step computation of M is simulated in $O(T(n))$ time. Hence, the resulting one-way-tape machine computes the same function as M . ■

What robustness does and does not say

These simulations establish that the models compute the same functions. They do not say that every choice of model gives exactly the same running time. For example, a direct three-tape

palindrome algorithm runs in linear time, while some restricted single-tape implementations may require substantially more time.

Reasonable variants of the Turing-machine model agree on computability and usually agree on broad complexity distinctions such as polynomial versus exponential time. This robustness is one reason Turing machines are useful as a foundation for complexity theory.

1.4 Universal Turing Machines

1.4.1 Machines as Data

A Turing machine has a finite description: it has a finite alphabet, finitely many states, and a finite transition table.² Therefore, the description of a machine can itself be encoded as a binary string. This simple observation allows one machine to receive another machine as input and simulate it.

Let M be a Turing machine. We write

$$\langle M \rangle \in \{0, 1\}^*$$

for a binary encoding of M . The encoding contains enough information to recover:

- the tape alphabet and an encoding of each tape symbol;
- the states, including the start and halting states;
- the number and roles of the tapes; and
- every entry of the transition function.

Because all these sets are finite, such an encoding can be constructed using ordinary binary representations and delimiters.

Definition 1.4.1 (Machine Represented by a String). *For a binary string α , let M_α denote the Turing machine represented by α . Strings that are not valid encodings may be assigned to a fixed default machine, such as a machine that immediately outputs 0 and halts.* $\diamond$

This convention ensures that every binary string denotes some machine. It is also convenient to permit harmless padding, so one machine may have more than one valid encoding.

Remark 1.4.1 (Programs as Data). *The distinction between program and data is not absolute. A machine description is data when it is read by another machine, but it specifies a program when interpreted as a transition table. This principle underlies interpreters, compilers, stored-program computers, and many later arguments in complexity theory.*

1.4.2 Definition and Basic Simulation

Definition 1.4.2 (Universal Turing Machine). *A universal Turing machine is a fixed Turing machine U such that, for every machine description α and every input x ,*

$$U(\alpha, x) = M_\alpha(x), \tag{1.5}$$

whenever $M_\alpha(x)$ halts. If $M_\alpha(x)$ does not halt, then the simulation by U also need not halt. $\diamond$

²Note that the infinite tape is not part of the *description*; only the sentence that “a Turing machine has an infinite tape” is part of the description; this sentence has finite number of words, and thus can be encoded as a finite string.

The notation $U(\alpha, x)$ means that U receives an encoding of the pair (α, x) . A standard delimiter or self-delimiting encoding can be used to distinguish the two components.

The universal machine maintains encoded versions of:

- the simulated machine's transition table;
- the simulated machine's current state;
- the contents of its tapes; and
- the positions of its tape heads.

To simulate one transition, U reads the simulated state and scanned symbols, looks up the appropriate transition in the encoded table, and updates the simulated configuration.

Analogy: Virtual Machines

A useful real-world analogy is a *virtual machine*. A virtual machine may behave like a complete computer, with its own processor, memory, disks, and operating system. Nevertheless, everything needed to specify it—its configuration, virtual hardware, and software—can be stored as finite data on a physical computer.

A physical computer can read this description and simulate the virtual computer step by step. The same physical computer can simulate many different virtual machines; which machine it simulates depends on the description it is given. Similarly, a universal Turing machine can take an encoding $\langle M \rangle$ as input and simulate the particular Turing machine M described by that encoding.

This analogy also illustrates the idea of *programs as data*. To the host computer, a virtual-machine image is simply a file containing bits. When those bits are interpreted by virtualization software, however, they describe a computer and determine its behavior. In the same way, $\langle M \rangle$ is an ordinary binary string when treated as input, but it becomes a program when interpreted as the description of a Turing machine.

1.4.3 Efficient Universal Simulation

The existence of a universal Turing machine does not immediately imply that the simulation is efficient. In this subsection, we discuss the accurate cost of simulating a machine M_α on input x using a universal machine U . We will first show that a naive simulation may introduce quadratic overhead, and then we will present a more efficient simulation that incurs only logarithmic overhead.

Theorem 1.4.1 (Quadratic Universal Turing Machine). *There exists a fixed Turing machine U such that, for every $\alpha, x \in \{0, 1\}^*$,*

$$U(\alpha, x) = M_\alpha(x)$$

whenever $M_\alpha(x)$ halts.

Furthermore, if M_α halts on x within T steps, then U halts within

$$C_\alpha T^2 \tag{1.6}$$

steps, where C_α depends on the fixed machine M_α , including its number of states, alphabet size, and number of tapes, but does not depend on x or T . $\diamond$

Before proving the theorem, let us emphasize its main strength: the input machine M_α is completely arbitrary. For example, M_α may have any fixed number of tapes and may use any

finite tape alphabet. In contrast, the universal machine U is fixed once and for all; indeed, U may be taken to be a standard three-tape Turing machine with binary tape alphabet. Nevertheless, U can simulate every machine M_α with at most quadratic overhead in its running time. Thus, although the constants involved may depend on the description α , the quadratic form of the bound is uniform over all input machines.

Proof of Theorem 1.4.1. We construct the required universal Turing machine U . In fact, U will be a standard three-tape Turing machine with binary tape alphabet $\{0, 1\}$.




Figure 1.4: A fixed universal Turing machine U simulating an arbitrary machine M_α . The machine U is a standard three-tape Turing machine with binary tape alphabet, whereas M_α may have any fixed number of tapes and any finite alphabet. Adapted from [AB09, Figure 1.6].

The basic idea is straightforward. Given the description α of a Turing machine M_α and an input x , the machine U maintains an encoding of the current configuration of M_α . It then repeatedly consults α to determine the next transition and updates the encoded configuration accordingly. This process is illustrated in Figure 1.4, where the “Description of M” and “Current state of M” segments can be placed to the beginning part of the work tape (right before the “Simulation of M’s work tape” segment), occupying only a constant number of cells.

The main point is that this simulation remains efficient even though M_α may have multiple tapes and an arbitrary finite alphabet. This follows from two observations:

1. Let Γ be the tape alphabet of M_α . Each symbol of Γ can be represented by a binary block of length

$$\lceil \log_2 |\Gamma| \rceil.$$

Since Γ is fixed once M_α is fixed, operations on these blocks incur only constant overhead relative to the running time of M_α .

2. Suppose that M_α has k tapes. Its k tapes can be encoded on a single tape by placing their contents on separate tracks and marking the locations of their heads. As in the proof of Lemma 1.3.2, one transition of M_α can be simulated by scanning the portion of the encoded tape visited so far. After t steps, this portion has length $O(t)$, so simulating the first T transitions requires

$$\sum_{t=1}^T O(t) = O(T^2)$$

time. Although k may be arbitrarily large across different machines, it is a constant for each fixed machine M_α .

The universal machine U performs these operations directly from the binary description α . Consequently, if M_α halts on x within T steps, then U simulates the computation and produces the same output in $O(T^2)$ time, with constants that may depend on α .

This finishes the proof of Theorem 1.4.1. ■

By a more involved construction, one can achieve a more efficient simulation of a universal Turing machine. The following [Theorem 1.4.2](#) shows that the overhead can be reduced to $O(T \log T)$ time. However, the proof of this theorem is much more technical and out of the scope of this undergraduate-level course. We present the proof in [Sections 1.4.4](#) and [1.4.5](#) as optional readings for interested students.

Theorem 1.4.2 (Efficient Universal Turing Machine). *There exists a fixed Turing machine U such that, for every $\alpha, x \in \{0, 1\}^*$,*

$$U(\alpha, x) = M_\alpha(x)$$

whenever $M_\alpha(x)$ halts.

Furthermore, if M_α halts on x within $T \geq 2$ steps, then U halts within

$$C_\alpha T \log T \tag{1.7}$$

steps, where C_α depends on the fixed machine M_α , including its number of states, alphabet size, and number of tapes, but does not depend on x or T . $\diamond$

1.4.4 *Proof of [Theorem 1.4.2](#) (Part 1/2)

Nothing from [Section 1.4.4](#) will be tested on quizzes or exams!

Before studying the $O(T \log T)$ -time universal Turing machine simulation, we examine two simpler examples involving powers of two. They illustrate two related but distinct ideas:

- Organizing a computation into exponentially increasing scales produces only logarithmically many levels.
- An operation can be expensive occasionally without making a long sequence of operations expensive on average.

Repeated squaring illustrates the first idea. A binary counter illustrates the second.

These ideas will become crucial in the real proof as we will show in [Section 1.4.5](#), where we use them in a more complicated way than in these examples. Thus, we believe it is helpful to understand them in a simpler context first, as we will do now in the current subsection.

1.4.4.1 Example 1: Computing powers by repeated squaring

Suppose we want to compute a^n , where n is a nonnegative integer. Multiplying a by itself repeatedly takes $n - 1$ multiplications when $n \geq 1$. Can we use fewer multiplications?

The key identities are

$$a^{2m} = (a^m)^2, \quad a^{2m+1} = a(a^m)^2.$$

Thus, to compute a^n , we can first compute $a^{\lfloor n/2 \rfloor}$, square the result, and multiply by a once more if n is odd.

More explicitly, the following recursive procedure computes a^n :

1. If $n = 0$, return 1.
2. If $n = 1$, return a .
3. Recursively compute

$$y = a^{\lfloor n/2 \rfloor}.$$

4. If n is even, return $y \cdot y$. If n is odd, return $a \cdot (y \cdot y)$.

The important point is that we compute y only once and then reuse it.

For example, to compute a^{13} , we can calculate

$$a^2 = a \cdot a, \quad a^3 = a \cdot a^2, \quad a^6 = a^3 \cdot a^3, \quad a^{12} = a^6 \cdot a^6, \quad a^{13} = a \cdot a^{12}.$$

This uses five multiplications rather than twelve.

Why the number of multiplications is logarithmic. Each recursive call roughly halves the exponent:

$$n, \quad \left\lfloor \frac{n}{2} \right\rfloor, \quad \left\lfloor \frac{n}{4} \right\rfloor, \quad \dots, \quad 1.$$

For $n \geq 1$, there are $\lfloor \log_2 n \rfloor$ reductions before reaching 1, and each reduction contributes at most two multiplications. Hence the procedure uses

$$O(\log(n+1))$$

multiplications.

An equivalent way to understand the method is through the binary representation of the exponent. For instance,

$$13 = 8 + 4 + 1, \quad \text{so} \quad a^{13} = a^8 a^4 a.$$

Repeated squaring constructs

$$a, \quad a^2, \quad a^4, \quad a^8, \quad \dots,$$

and the binary digits of n tell us which powers to multiply together.

A warning about the cost model. This is a bound on the *number of multiplications*, not necessarily on the number of Turing machine steps. Multiplication is not a constant-time operation when its operands are arbitrarily long integers. For example, if $a \geq 2$ is a fixed integer, the binary representation of a^n already has $\Theta(n)$ bits.

The lesson we want here is structural:

Doubling the scale at each level gives logarithmically many levels.

1.4.4.2 Example 2: Incrementing a binary counter

Consider a binary counter initially equal to 0. We increment it T times, where $T \geq 1$. Allocate

$$b = \lceil \log_2(T+1) \rceil$$

bits, which is enough to represent every value from 0 through T .

To increment the counter, start at the least significant bit:

1. Change each consecutive 1 into a 0, carrying to the next bit.
2. Change the first 0 into a 1, and stop.

For example,

$$0111 \longrightarrow 1000$$

changes four bits, whereas

$$1000 \longrightarrow 1001$$

changes only one.

We measure the cost of an increment by the number of bits it changes.

A worst-case bound is not the whole story. One increment can change all $b = O(\log(T + 1))$ bits. Multiplying this worst-case cost by the number of increments gives the valid upper bound

$$O(T \log(T + 1)).$$

However, this overestimates the total cost: increments with long carry chains are rare.

To obtain a sharper bound, we count how often each bit changes over the entire sequence.

Number the bits $0, 1, \dots, b - 1$, starting with the least significant bit. Then:

- Bit 0 changes on every increment.
- Bit 1 changes on every second increment.
- Bit 2 changes on every fourth increment.
- In general, bit i changes on every 2^i -th increment.

Consequently, the total number of bit changes is

$$\begin{aligned} \sum_{i=0}^{b-1} \left\lfloor \frac{T}{2^i} \right\rfloor &\leq T \sum_{i=0}^{b-1} \frac{1}{2^i} \\ &< T \sum_{i=0}^{\infty} \frac{1}{2^i} \\ &= 2T. \end{aligned}$$

Thus all T increments together cost $O(T)$, even though a single increment can cost $\Theta(\log(T + 1))$.

What amortized analysis means. We say that incrementing the counter has $O(1)$ *amortized cost*: the total cost of T increments, starting from zero, is $O(T)$.

This is not a probabilistic statement. We are not assuming that the counter values are random, nor that every increment is cheap. Instead, we prove a bound on the total cost of the entire sequence.

The useful accounting principle is

$$\text{Total cost} = \sum_{\text{levels}} (\text{cost per event}) (\text{number of events}).$$

For the counter, a level- i event is a change to bit i . It costs 1 and occurs at most $T/2^i$ times.

Connecting these examples to tape simulation. As we will see in [Section 1.4.5](#), the universal Turing machine simulation uses buffers whose sizes grow geometrically:

$$2, \quad 4, \quad 8, \quad 16, \quad \dots$$

As in repeated squaring, there are only $O(\log(T + 1))$ relevant levels during a computation of T steps.

As with a binary counter, activity at high levels is infrequent. The buffer construction will ensure that reorganizing level i costs $O(2^i)$, but occurs only $O(T/2^i)$ times. Therefore, each level contributes at most

$$O(2^i) \cdot O(T/2^i) = O(T)$$

work. Summing over all relevant levels gives

$$O(T) \cdot O(\log(T + 1)) = O(T \log(T + 1)).$$

Notice the difference from the binary counter: changing a high-order bit costs only 1, whereas reorganizing a large buffer costs proportionally to its size. This is why the counter has $O(T)$ total cost, while the buffer simulation has an $O(T \log(T + 1))$ bound.

The proof that follows makes this accounting argument precise by constructing the buffers and showing why a large rearrangement cannot happen again until sufficiently many simulated head moves have occurred.

1.4.5 *Proof of Theorem 1.4.2 (Part 2/2)

Nothing from Section 1.4.5 will be tested on quizzes or exams!

We now proceed to the real proof of Theorem 1.4.2.

We use the input/output-tape convention of the theorem: U simulates M 's input and output tapes directly. The main task is to simulate M 's k work tapes using a fixed number of work tapes.

1.4.5.1 The main idea: move the contents, not the simulated heads

Let Γ be the work-tape alphabet of M . We view one work tape of U as having k parallel *tracks*. A cell of this tape stores one symbol for each track, so it represents a tuple of k symbols.

These tracks are not independent physical tapes: they share a single physical head. Nevertheless, we will arrange that, between simulated steps, the symbol under each of M 's work-tape heads is stored at position 0 of its corresponding track.

Thus, when a simulated head moves, we move the represented tape contents relative to position 0:

- A simulated move to the right brings the next symbol on the right into position 0.
- A simulated move to the left brings the next symbol on the left into position 0.

If we physically shifted an entire track after every move, one simulated step could cost $O(T)$, giving an $O(T^2)$ simulation. Instead, we introduce unused space into each track and use it as a collection of buffers.

Implementation convention. For readability, we describe cells as storing tuples from an alphabet that depends on M . The actual universal machine has a fixed alphabet: it encodes each such tuple by a block of fixed-alphabet symbols. For a fixed M , the block length is constant, so this changes the running time by only a factor depending on M .

1.4.5.2 Buffers and their meaning

Introduce a padding symbol $\diamond$, distinct from every symbol of Γ , including the ordinary tape blank.

The distinction is important:

- An ordinary blank represents an actual cell of M 's tape.
- A $\diamond$ represents unused storage and is ignored.

For example, the strings

$$010, \quad 0 \diamond 10, \quad 0 \diamond \diamond 1 \diamond 0$$

represent the same sequence of simulated tape symbols.

We use a two-sided physical tape, with positions indexed by the integers. Position 0 is the *home position*. For each $i \geq 1$, define a right zone and a left zone by

$$R_i = \{2^i - 1, \dots, 2^{i+1} - 2\}, \quad L_i = \{-p : p \in R_i\}.$$

Both zones have 2^i cells. The layout is therefore

$$\cdots L_3 L_2 L_1 \boxed{0} R_1 R_2 R_3 \cdots,$$

where the zone sizes, moving outward from the center, are

$$2, 4, 8, 16, \dots$$

Starting the indexing at 1 ensures that every zone has an integer half-capacity.

For each track separately, we maintain these invariants:

1. Position 0 contains the symbol currently under the corresponding head of M .
2. Reading the track and ignoring all $\diamond$'s gives the simulated tape contents, in the correct order.
3. Each zone of size 2^i is either empty, half full, or full of non-padding symbols.
4. The total number of non-padding symbols in L_i and R_i is exactly 2^i .

The last two conditions mean that the occupancies of L_i and R_i must be one of

$$(0, 2^i), \quad (2^{i-1}, 2^{i-1}), \quad (2^i, 0).$$

Initially, every zone is conceptually half full of ordinary blank symbols, and position 0 also contains a blank. There is no need to initialize infinitely many cells: a zone is initialized only when it is first reached. Beyond the initialized region, all simulated symbols are still blank.

1.4.5.3 Simulating one transition

At the beginning of a simulated step, U 's physical head is at position 0. It reads the tuple of current work-tape symbols, consults M 's transition table, and writes the updated symbols at position 0.

It then handles the prescribed head movement on each track separately. A stationary head requires no rearrangement. We explain a move to the right (see also [Figure 1.5](#) for an illustration); a move to the left is symmetric.

1.4.5.4 Simulating a move to the right

A move to the right must bring the nearest represented symbol on the right into position 0, while placing the old position-0 symbol on the left.

Find the smallest i such that R_i is nonempty. Then

$$R_1, \dots, R_{i-1} \text{ are empty,}$$

and, by the paired-occupancy invariant,

$$L_1, \dots, L_{i-1} \text{ are full.}$$

Write

$$h = 2^{i-1}.$$




Figure 1.5: A rightward head move on one track, assuming the transition does not change the symbol e . The simulator brings f to position 0, places g in R_1 , moves c, d into L_2 , and leaves e in L_1 . The symbol ◇ denotes unused storage, not an ordinary tape blank. Only zones through level 2 are rearranged; all farther zones are unchanged.

The total capacity of the smaller zones on either side is

$$\sum_{j=1}^{i-1} 2^j = 2^i - 2 = 2h - 2,$$

and their total half-capacity is

$$\sum_{j=1}^{i-1} 2^{j-1} = h - 1.$$

We now rearrange only

$$L_i, \dots, L_1, 0, R_1, \dots, R_i.$$

Right side. Take the first h non-padding symbols of R_i , in order starting from the end nearest position 0.

- Put the first symbol at position 0.
- Distribute the remaining $h - 1$ symbols among $R_1, \dots, R_{i-1}$, making each of these zones half full.

Preserve the order of the represented symbols. The occupancy of R_i decreases by h .

Left side. The smaller left zones contain $2h - 2$ symbols. Together with the old symbol at position 0, there are $2h - 1$ symbols to place on the left.

- Move the h symbols farthest from the new head position into L_i , placing them after its existing symbols in left-to-right order.
- Distribute the remaining $h - 1$ symbols among $L_1, \dots, L_{i-1}$, making each of these zones half full.

Again, preserve the order of all represented symbols. The occupancy of L_i increases by h .

Why the invariants are preserved. Before the rearrangement, the occupancies of (L_i, R_i) were either

$$(h, h) \quad \text{or} \quad (0, 2h),$$

because R_i was nonempty. Afterward, they are respectively

$$(2h, 0) \quad \text{or} \quad (h, h).$$

Thus the allowed occupancies and their sum are preserved.

All smaller zones are now half full. Position 0 contains the next simulated tape symbol, and the old position-0 symbol is immediately to its left in the represented sequence. Consequently, this rearrangement correctly simulates one rightward head move.

For $i = 1$, there are no smaller zones. The operation simply transfers the old central symbol into L_1 and the nearest symbol of R_1 into the center.

1.4.5.5 The cost of one rearrangement

Call an operation whose outermost affected zones are L_i, R_i a *level- i rearrangement*.

Its entire affected region has length

$$1 + 2 \sum_{j=1}^i 2^j = O(2^i).$$

Using a scratch tape, U can copy the represented symbols in order and redistribute them with a constant number of scans of this region. Zone boundaries and half-capacities can be marked on auxiliary tracks.

Therefore, a level- i rearrangement costs

$$O_M(2^i).$$

This includes returning the physical head to position 0. Initializing a newly reached zone also fits within this bound.

The crucial point is that expensive rearrangements occur rarely.

1.4.5.6 Why large rearrangements are rare

Fix one track and a level i . Consider any rearrangement that reaches level i or farther.

Immediately afterward, every smaller zone

$$L_1, R_1, \dots, L_{i-1}, R_{i-1}$$

is half full. Hence the smaller zones contain exactly

$$s_i = \sum_{j=1}^{i-1} 2^{j-1} = 2^{i-1} - 1$$

represented symbols on each side of position 0.

Until another rearrangement reaches level i , each rightward move decreases the number of symbols in the smaller right zones by one, and each leftward move increases it by one. The analogous statement holds on the left.

To reach level i again, all smaller zones on one side must first be empty. Starting with s_i symbols on each side, this requires at least s_i moves, followed by the move that actually reaches

level i . Thus two rearrangements reaching level i or farther are separated by at least

$$s_i + 1 = 2^{i-1}$$

head moves.

The same lower bound applies before the first such rearrangement, because initially all zones are half full.

Since a work-tape head makes at most T moves during T steps, the number of rearrangements reaching level i or farther is at most

$$\left\lfloor \frac{T}{2^{i-1}} \right\rfloor.$$

In particular, this also bounds the number of level- i rearrangements.

Moreover, a level can be reached only if $2^{i-1} \leq T$. Thus the largest possible level is

$$1 + \lfloor \log_2 T \rfloor.$$

1.4.5.7 Summing the costs

For one track, the total rearrangement time is at most

$$\begin{aligned} \sum_{i=1}^{1+\lfloor \log_2 T \rfloor} O_M(2^i) \left\lfloor \frac{T}{2^{i-1}} \right\rfloor &\leq O_M \left(\sum_{i=1}^{1+\lfloor \log_2 T \rfloor} 2^i \frac{T}{2^{i-1}} \right) \\ &= O_M \left(\sum_{i=1}^{1+\lfloor \log_2 T \rfloor} 2T \right) \\ &= O_M(T \log(T+1)). \end{aligned}$$

There are k tracks, where k depends only on M , so their combined cost has the same asymptotic bound. Reading the transition table, updating the state, and handling the input and output tapes add only $O_M(T)$ time.

Therefore U correctly simulates M within

$$O_M(T \log(T+1))$$

steps.

The idea to remember. A buffer at level i costs $O(2^i)$ to reorganize, but it is involved only once per $\Omega(2^i)$ simulated head moves. Each level therefore contributes only $O(T)$ total work. There are $O(\log T)$ levels, giving $O(T \log T)$ time overall.

This finishes the proof of [Theorem 1.4.2](#).

Takeaway

A single fixed machine can simulate every other Turing machine with only modest overhead. Thus, computational universality does not require building a different physical device for every algorithm; the algorithm itself can be supplied as encoded data.

1.5 The Church–Turing Thesis

We have shown that several formal models can simulate one another, but a broader question remains: does the Turing-machine model capture every computation that could be carried out

by any effective mechanical procedure? The Church–Turing thesis proposes that it does.

1.5.1 The Standard Church–Turing Thesis

Church–Turing thesis: Every function that can be computed by an effective mechanical procedure can be computed by a Turing machine.

The phrase *effective mechanical procedure* refers informally to a finite collection of unambiguous rules that can be followed step by step without requiring creativity, insight, or infinitely precise information.

The Church–Turing thesis is not a theorem. One side of the thesis refers to the informal notion of an effective procedure, so the statement cannot be proved using only formal mathematics. It is a hypothesis connecting an informal concept to a formal model.

Evidence for the thesis includes the following:

- Many independently proposed formal models of computation, including Turing machines, lambda calculus, and general recursive functions, define the same class of computable functions.
- Standard programming languages can simulate Turing machines, and Turing machines can simulate programs written in standard programming languages.
- Natural changes to the details of the Turing-machine model do not change the class of computable functions.
- No generally accepted effective procedure has been discovered that computes a function uncomputable by a Turing machine.

Implications for the limits of computation

If we accept the standard Church–Turing thesis, then a proof that no Turing machine computes a particular function establishes a fundamental limit. The conclusion is not merely that present computers lack sufficient memory or that current programmers have not found the correct algorithm. Rather, no algorithmic procedure can compute that function on every input.

This distinction will become important when we study undecidability. Some problems are difficult because known algorithms use many resources. Other problems are uncomputable: no Turing-machine algorithm exists at all.

1.5.2 The Extended Church–Turing Thesis

The standard thesis concerns what is computable, without requiring the simulation to be efficient. Complexity theory motivates a stronger statement.

Extended Church–Turing thesis: Every computation carried out by a physically realizable computational device can be simulated efficiently, usually meaning with at most polynomial overhead, by a probabilistic Turing machine.

The addition of “efficiently” is crucial. An exponential-time simulation still supports the standard Church–Turing thesis, because the same function remains computable. It may nevertheless contradict the extended thesis because exponential overhead is not considered efficient.

Thesis	Main question	Required simulation
Standard	What can be computed?	Any finite overhead is allowed
Extended	What can be computed efficiently?	Polynomial overhead

Table 1.2: The standard and extended Church–Turing theses make different claims. The extended thesis is strictly stronger because it imposes an efficiency requirement.

1.5.3 Quantum Computing and the Extended Thesis

Quantum computers are physical models whose evolution is governed by quantum mechanics. They can be formalized by quantum circuits or quantum Turing machines.

A classical Turing machine can simulate any finite quantum computation to any required finite precision. The classical simulation may, however, require exponential time because a general state of m quantum bits is described using 2^m complex amplitudes.

Let:

- BPP denote the class of decision problems solvable in polynomial time by a “probabilistic” (we will discuss this in later chapters) classical Turing machine with bounded error; and
- BQP denote the class of decision problems solvable in polynomial time by a quantum computer with bounded error.

If quantum computers can efficiently solve a problem that no probabilistic classical Turing machine can solve efficiently, then

$$\text{BQP} \not\subseteq \text{BPP}.$$

Such a separation would contradict the classical probabilistic version of the extended Church–Turing thesis.

It has not been proved that $\text{BQP} \not\subseteq \text{BPP}$. Quantum algorithms provide strong evidence that quantum computation may violate the extended Church–Turing thesis, but this conclusion remains conditional on unresolved complexity-theoretic questions.

Quantum computing does not, according to the standard formal models, violate the standard Church–Turing thesis. A quantum computation can be simulated by a classical Turing machine, although possibly with exponential overhead. It changes our understanding of efficient computation, not the class of functions that are computable in principle.

Remark 1.5.1 (What Quantum Computing Challenges). *Quantum computing challenges the claim that every reasonable physical computation has an efficient classical simulation. It does not provide a method for solving Turing-uncomputable problems. Thus, it is evidence against some formulations of the extended thesis, not against the standard thesis.*

1.5.4 The Turing Machine as a Final Model

If we accept the standard Church–Turing thesis, then introducing a more realistic programming language, a different processor, more tapes, random-access memory, parallel hardware, or quantum mechanics does not enlarge the class of computable functions beyond Turing computability.

This does not mean that all models are equally convenient or equally efficient:

- Different programming languages may make different algorithms easier to express.
- Different machine models may change running time by polynomial or even exponential factors.

- Randomness, parallelism, and quantum computation may change which problems can be solved efficiently.
- Restricted models such as finite automata remain valuable when limited memory is part of the problem being studied.

The main point is more fundamental: under the standard Church–Turing thesis, changing to another physically realizable model does not allow us to compute a genuinely uncomputable function.

If the standard Church-Turing thesis is correct, then the Turing machine is an eventual final model of computability: future models may alter convenience or efficiency, but they will not enlarge the set of functions that are computable by effective procedures.

1.6 Closing Remarks

We emphasize that Turing machines are not the only mathematical model of computation. Other models include the lambda calculus, recursive functions, register machines, and various idealized programming languages. Although these models describe computation in different ways, they are *equivalent* at the level of computability: each can compute exactly the functions computable by a Turing machine. Thus, the Turing machine should not be regarded as the uniquely correct description of computation, but as one particularly useful formalization among several equivalent ones.

We focus on Turing machines because their operation is concrete and easy to visualize: a machine reads and writes symbols, moves its tape heads, and changes its finite state. These elementary actions make it possible to describe algorithms precisely and to measure resources such as running time and memory usage. Moreover, the distinction between the finite control and the unbounded tape closely reflects the familiar distinction between a fixed program and its working memory. Turing machines therefore provide a simple, rigorous, and convenient foundation for studying computability and computational complexity, even though other equivalent models could serve the same fundamental purpose.

We also emphasize that we do not expect you to memorize every detail of the formal definition of a Turing machine or its transition rules. Rather, the purpose of studying Turing machines is to see how computation can be described with mathematical precision by reducing it to a small collection of elementary operations. Much like building with LEGO bricks, each natural step of an algorithm can be decomposed into finitely many simple Turing-machine steps. Despite its simplicity, this model has proved remarkably robust: every standard, reasonable model of general-purpose computation developed so far can compute the same class of functions as a Turing machine. This observation provides the main evidence for the Church-Turing thesis, which asserts that anything computable by an effective algorithm can be computed by a Turing machine. Thus, Turing machines give us a simple and implementation-independent foundation for reasoning rigorously about what can be computed and how efficiently it can be computed.

Updates upto here

Bibliography

- [AB09] Sanjeev Arora and Boaz Barak. *Computational complexity: a modern approach*. Cambridge University Press, 2009. [14](#), [18](#)