

CSCI3160 Design and Analysis of Algorithms (2025 Fall)

Approximation Algorithms 1: Vertex Cover and MAX-3SAT

Instructor: Xiao Liang¹

Department of Computer Science and Engineering
The Chinese University of Hong Kong

¹These slides are primarily based on materials prepared by Prof. Yufei Tao (please refer to [Prof. Tao's version from 2024 Fall](#) for the original content). Some modifications have been made to better align with this year's teaching progress, incorporating student feedback, in-class interactions, and my own teaching style and research perspective.

Motivation

We have learned several algorithms that help us solve problems efficiently:

- Sorting in time $O(n \log n)$
- Matrix multiplication in time $O(n^{2.81})$
- FFT-based polynomial multiplication in time $O(n \log n)$
- Activity selection in time $O(n \log n)$
- ...
- All-Pairs Shortest Paths in time $O(|V|(|V| + |E|) \log(|V|))$

However: there are still many problems **of practical significance** for which no efficient (i.e., polynomial-time) algorithms are currently known to us humans.

Example 1: Graph-3-Coloring

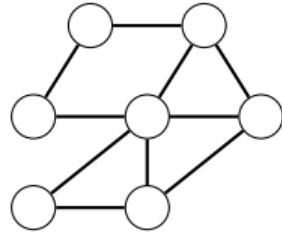
A **graph coloring** assigns colors to the vertices of a graph so that no two adjacent vertices share the same color.

In the **Graph-3-Coloring** problem, we ask:


Can the vertices of a given graph be colored with at most 3 colors such that adjacent vertices have different colors?

Extension: 3-coloring is a special case of the general k-coloring problem.

Exemplary Graphs



(a) A 3-colorable graph



(b) A 3-coloring scheme




Figure: A non-3-colorable graph

Application 1: Map Coloring

- Each region on a map can be represented as a vertex in a graph.
- An edge connects two vertices if the corresponding regions share a common border.
- The goal is to color each region so that no two adjacent regions have the same color.
- If the map can be colored with 3 colors, the corresponding graph is 3-colorable.
- This has applications in geography, political boundary planning, and resource distribution.

Application 2: Scheduling Exams

- Vertices represent courses.
- An edge connects two courses if they have students in common.
- Goal: assign each course to one of 3 time slots, avoiding conflicts.
- 3-coloring determines if this is possible.
- If there are k courses under consideration, then it corresponds to the k -coloring problem.

Application 3: Frequency Assignment

- Nodes represent transmitters in a communication network.
- Edges represent interference (proximity).
- Assign 3 frequencies to avoid interference between neighbors.
- 3-coloring tells us if this can be done using only 3 frequencies.
- In general k -coloring tells us if this can be done using only k frequencies.

Subset Sum Problem and Applications

Problem Statement:

- Given a set of integers $S = \{x_1, x_2, \dots, x_n\}$ and a target integer T ,
- Does there exist a subset $S' \subseteq S$ such that the sum of elements in S' is exactly T ?

Example 1:

- $S = \{3, 34, 4, 12, 5, 2\}$, $T = 9$
- Yes: $\{4, 5\}$ or $\{3, 4, 2\}$ sum to 9.

Example 2:

- $S = \{3, 5, 9, 13\}$, $T = 7$
- No!

Application 1: Budget Allocation

Scenario:

- A company has a list of proposed projects, each with a known cost.
- The total available budget is a fixed amount T .
- The goal is to determine if there is a combination of projects whose total cost exactly matches the budget.

Example:

- Projects: $\{P_1 : \$30k, P_2 : \$50k, P_3 : \$20k, P_4 : \$40k\}$
- Budget: $\$90k$
- Is there a subset of projects that costs exactly $\$90k$?
- Yes: $\{P_2, P_3, P_4\} \rightarrow 50k + 20k + 20k = 90k$

Relevance:

- Subset sum helps in decision support for finance and planning.
- Used in automated budget optimization tools and resource allocation systems.

Application 2: Packing and Logistics

Scenario:

- A shipping company needs to fill containers with items of different weights or volumes.
- Goal: Select a subset of items that perfectly fills a container of limited capacity.

Example:

- Items: $\{w_1 = 3\text{kg}, w_2 = 7\text{kg}, w_3 = 2\text{kg}, w_4 = 6\text{kg}\}$
- Container capacity: 9kg
- Feasible subset: $\{w_2, w_3\} \rightarrow 7 + 2 = 9$

Relevance:

- Subset sum models the core problem in bin packing and cargo loading.
- Used in warehouse automation, shipping logistics, and supply chain optimization.

Unfortunately, we currently do not know any polynomial-time algorithms for Graph 3-Coloring, Subset Sum, or many other problems that have broad applications and significant real-world importance.

A central research theme for Theoretical Computer Science (TCS):

- **What can we do about these hard problems?**

Two branches of computer science have been developed centered around this theme:

- **Computation Complexity:** Seeks to understand the inherent difficulty of problems and classify them based on resource requirements. (Not our focus in this course.)
- **Approximation Algorithms:** Develops efficient algorithms that produce near-optimal solutions to hard problems. (This will be our focus for the remaining lectures.)

Branch 1: Computational Complexity

Goal: Understand the fundamental limits of computation.

- Classifies problems into complexity classes such as:
 - **P:** Problems solvable in polynomial time.
 - **NP:** Problems whose solutions can be verified in polynomial time.
 - **NP-complete:** The hardest problems in NP — if any one of them can be solved in polynomial time, then all of NP can.
- Addresses profound open questions like:
 - **P vs NP:** Can every efficiently verifiable problem also be efficiently solvable?
- Also studies:
 - Reductions between problems (to compare difficulty)
 - Space and time trade-offs
 - Randomized and quantum complexity classes

Takeaway: Computational complexity helps us understand *why* certain problems are hard, and how that hardness is structured.

Branch 2: Approximation Algorithms

Goal: Design efficient algorithms that find near-optimal solutions for hard optimization problems.

- When exact solutions are computationally infeasible (e.g., NP-hard problems), we aim for **good enough** solutions in **polynomial time**.
- An approximation algorithm returns a solution whose value is within a provable factor of the optimum.
- Central questions in this area:
 - How close can we get to the optimal solution in polynomial time?
 - Are there limits (hardness of approximation) beyond which we can't do better unless $P = NP$?
- Techniques include:
 - Greedy methods
 - LP/SDP relaxations
 - Randomized rounding

Takeaway: Approximation algorithms offer a practical path forward for solving hard problems when exact solutions are out of reach.

An introductory Journey to Approximation Algorithms

For a rigorous discussion about Approximation Algorithms, we first need to borrow some concepts from Computational Complexity.

What's the exact meaning of "no polynomial-time algorithms are known" for some problem (e.g., Graph 3-Coloring, Subset Sum)?

- We must be precise about the **computation model** we are using.

Long story short, **Turing Machine** has become the standard model of computation in complexity theory:

- We won't explain what a Turing Machine is. But in short, it is a slightly more powerful model than the RAM model we utilize in this course so far.
- It is mathematically simple yet powerful enough to simulate any "reasonable" algorithm.
- It is **robust**: polynomial-time computations on one reasonable model can be simulated in polynomial time on a Turing Machine.
- It is **closed under composition**: combining polynomial-time Turing Machines results in a polynomial-time Turing Machine.

Church-Turing Thesis

Informal Statement:

Any function that can be computed by a "reasonable" mechanical procedure (i.e., algorithm) can be computed by a Turing Machine.

Key Points:

- It is a **foundational hypothesis** in computer science — not a formal theorem.
- Supported by the equivalence of many independent computational models:
 - Turing Machines
 - Lambda calculus
 - Recursive functions
 - RAM machines
- All known models of computation that align with our intuitive notion of an algorithm are **no stronger than** in power.
- Thus, this thesis conjecture that the Turing Machine is an adequate model for defining what is computable (or an Algorithm).

A Brief History of the Church-Turing Thesis

Early 1900s: Mathematicians, including **David Hilbert**, were seeking to formalize all of mathematics.

- Hilbert posed the famous **Entscheidungsproblem** (decision problem):
 - *Is there a general algorithm to determine whether a given mathematical statement is provable?*
- This led to the deeper question: **What exactly is an algorithm or effective procedure?**

Alonzo Church (1936):

- Proposed the **lambda calculus** as a formal model of computation.
- Introduced the notion of **computable functions**.
- Argued that this model captured all effectively calculable functions.


Alan Turing (1936):

- Independently tackled the same question.
- Introduced the **Turing Machine** in his groundbreaking paper:
 - “On Computable Numbers, with an Application to the Entscheidungsproblem”
- Showed that Turing Machines could simulate any mechanical computation.

Fun fact: Church was Turing's academic advisor at Princeton after this work!



(a) Alan Turing



(b) Alonzo Church

In computer science, there is a set of **NP-hard** problems such that

- nobody has found a polynomial-time algorithm for **any** of those problems;
- no polynomial-time algorithms can exist for **any** of those problems **unless** $\mathcal{P} = \mathcal{NP}$.

- \mathcal{P} = the set of problems that can be solved in polynomial time on a Turing machine
- \mathcal{NP} = the set of problems whose solution can be efficiently verified by a Turing machine

Turing machines are formalized in CSCI3130 (Formal Languages and Automata Theory), and so is the notion of NP-hard.

Whether $\mathcal{P} = \mathcal{NP}$ is still unsolved to this day.

What can we do if a problem is NP-hard?

The rest of the course will focus on a principled approach for tackling NP-hard problems: **approximation**.

In many problems, even though an optimal solution may be expensive to find, we can find **near-optimal** solutions efficiently.

Next, we will see two examples: **vertex cover** and **MAX-3SAT**.


The Vertex Cover Problem

$G = (V, E)$ is a simple undirected graph.

A subset $S \subseteq V$ is a **vertex cover** of G if every edge $\{u, v\} \in E$ is incident to at least one vertex in S .

The V.C. Problem: Find a vertex cover of the smallest size.

Example:



An optimal solution is $\{a, f, c, e\}$.

The vertex cover problem is NP-hard.

- No one has found an algorithm solving the problem in time polynomial in $|V|$.
- Such algorithms cannot exist if $\mathcal{P} \neq \mathcal{NP}$.

Approximation Algorithms

\mathcal{A} = an algorithm that, given any legal input $G = (V, E)$, returns a vertex cover of G .

OPT_G = the smallest size of all the vertex covers of G .

\mathcal{A} is a **ρ -approximate algorithm** for the vertex cover problem if, for any legal input $G = (V, E)$, \mathcal{A} can return a vertex cover with size at most $\rho \cdot OPT_G$.

The value ρ is the **approximation ratio**.

We say that \mathcal{A} achieves an approximation ratio of ρ .

Consider the following algorithm.

Input: $G = (V, E)$

$S = \emptyset$

while E is not empty **do**

 pick an arbitrary edge $\{u, v\}$ in E

 add u, v to S

 remove from E all the edges of u and all the edges of v


return S

It is easy to show:

- S is a vertex cover of G ;
- The algorithm runs in time polynomial to $|V|$ and $|E|$.

We will prove later that the algorithm is 2-approximate.

Example:



Suppose we start by picking edge $\{b, c\}$.

Then, $S = \{b, c\}$ and $E = \{\{a, e\}, \{a, d\}, \{d, e\}, \{d, f\}\}$.

Any edge in E can then be chosen. Suppose we pick $\{a, e\}$.

Then, $S = \{a, b, c, e\}$ and $E = \{\{d, f\}\}$.

Finally, pick $\{d, f\}$.

$S = \{a, b, c, d, e, f\}$ and $E = \emptyset$.

Theorem 1: The algorithm returns a set of at most $2 \cdot OPT_G$ vertices.

Let M be the set of edges picked.

Example: In the previous example, $M = \{\{b, c\}, \{a, e\}, \{d, f\}\}$.

Lemma 1: The edges in M do not share any vertices.

Proof: Suppose that M has edges e_1 and e_2 both incident to a vertex v . W.l.o.g., assume that e_1 was picked before e_2 . After picking e_1 , the algorithm deleted all the edges of v , because of which e_2 could not have been picked, giving a contradiction. \square

Lemma 2: $|M| \leq OPT_G$.

Proof: Any vertex cover must include at least one vertex of each edge in M . $|M| \leq OPT_G$ follows from Lemma 1. \square

Theorem 1 holds because the algorithm returns exactly $2|M|$ vertices.

The MAX-3SAT Problem

A **variable**: a boolean unknown x whose value is 0 or 1.

A **literal**: a variable x or its negation \bar{x} .

A **clause**: the OR of 3 literals with different variables.

S = a set of clauses

\mathcal{X} = the set of variables appearing in at least one clause of S

A **truth assignment** of S : a function from \mathcal{X} to $\{0, 1\}$.

A truth assignment f **satisfies** a clause in S if the clause evaluates to 1 under f .

The MAX-3SAT Problem: Let S be a set of n clauses. Find a truth assignment of S to maximize the number of clauses satisfied.

Example:

$$\begin{aligned} S = \{ & x_1 \vee x_2 \vee x_3, \\ & x_1 \vee x_2 \vee \bar{x}_3, \\ & x_1 \vee \bar{x}_2 \vee x_3, \\ & x_1 \vee \bar{x}_2 \vee \bar{x}_3, \\ & \bar{x}_1 \vee x_3 \vee x_4, \\ & \bar{x}_1 \vee x_3 \vee \bar{x}_4, \\ & \bar{x}_1 \vee \bar{x}_3 \vee x_4, \\ & \bar{x}_1 \vee \bar{x}_3 \vee \bar{x}_4 \}. \end{aligned}$$

$n = 8$ and $\mathcal{X} = \{x_1, x_2, x_3, x_4\}$.

The truth assignment $x_1 = x_2 = x_3 = x_4 = 1$ satisfies 7 clauses. It is impossible to satisfy 8.

The MAX-3SAT problem is NP-hard.

- No one has found an algorithm solving the problem in time polynomial in n .
- Such algorithms cannot exist if $\mathcal{P} \neq \mathcal{NP}$.

Approximation Algorithms

\mathcal{A} = an algorithm that, given any legal input S , returns a truth assignment of S .

OPT_S = the largest number of clauses that a truth assignment of S can satisfy.

Z_S = the number of clauses satisfied by the truth assignment \mathcal{A} returns.

- Z_S is a random variable if \mathcal{A} is randomized.

\mathcal{A} is a **randomized ρ -approximate algorithm** for MAX-3SAT if $E[Z_S] \geq \rho \cdot OPT_S$ holds for any legal input S .

The value ρ is the **approximation ratio**.

We also say that \mathcal{A} achieves an approximation ratio of ρ in expectation.

Consider the following algorithm.

Input: a set S of clauses with variable set \mathcal{X}

for each variable $x \in \mathcal{X}$ **do**

toss a fair coin

if the coin comes up heads **then** $x \leftarrow 1$

else $x \leftarrow 0$

It is clear that the algorithm runs in $O(n)$ time.

Next, we show that the algorithm achieves an approximation ratio $7/8$ in expectation.

Theorem 2: The algorithm produces a truth assignment that satisfies $\frac{7}{8}n$ clauses in expectation.

Proof: It suffices to show that each clause is satisfied with probability $7/8$. W.l.o.g., suppose that the clause is $x_1 \vee x_2 \vee x_3$. The clause is 0 if and only if x_1 , x_2 , and x_3 are all 0. The probability for $x_1 = x_2 = x_3 = 0$ is $1/8$. □

Think: What about a clause like $x_1 \vee x_2 \vee \bar{x}_3$?