

CSCI3160 Design and Analysis of Algorithms (2025 Fall)


Approximation Algorithms 4: k -Center

Instructor: Xiao Liang¹

Department of Computer Science and Engineering
The Chinese University of Hong Kong

¹These slides are primarily based on materials prepared by Prof. Yufei Tao (please refer to [Prof. Tao's version from 2024 Fall](#) for the original content). Some modifications have been made to better align with this year's teaching progress, incorporating student feedback, in-class interactions, and my own teaching style and research perspective.

Given 2D points p and q , we use $\text{dist}(p, q)$ to represent their Euclidean distance.



In this lecture, we will make the assumption that $\text{dist}(p, q)$ can be computed in polynomial time.

P = a set of n points in 2D space.

Given a point $p \in P$, define its distance to a subset $C \subseteq P$ as

$$\text{dist}_C(p) = \min_{c \in C} \text{dist}(p, c).$$

The **penalty** of C is

$$\text{pen}(C) = \max_{p \in P} \text{dist}_C(p).$$


The k -Center Problem: Find a subset $C \subseteq P$ with size $|C| = k$ that has the smallest penalty.

Example:

P = the set of black points

$k = 3$

$C = \{c_1, c_2, c_3\}$



Application 1: Emergency Facility Placement

Scenario: A city plans to build k emergency facilities (e.g., hospitals, fire stations).

Goal:

- Minimize the maximum response time to any area in the city.
- Ensure that every resident is as close as possible to at least one facility.

Model:

- Points: population centers or demand locations.
- Distance: travel time or road distance.

Application 2: Data Clustering

Scenario: In machine learning or data mining, you want to group data into k compact clusters.

Goal:

- Assign each data point to its nearest cluster center.
- Minimize the largest distance between any point and its assigned center.

Use Case:

- Prototype selection in large datasets.
- Reducing latency in content delivery networks.

The problem is NP-hard.

- No one has found an algorithm solving the problem in time polynomial in n and k .
- Such algorithms cannot exist if $\mathcal{P} \neq \mathcal{NP}$.

\mathcal{A} = an algorithm that, given any legal input P , returns a subset of P with size k .

Denote by OPT_P the smallest penalty of all subsets $C \subseteq P$ satisfying $|C| = k$.

\mathcal{A} is a **ρ -approximate algorithm** for the k -center problem if, for any legal input P , \mathcal{A} can return a set C with penalty at most $\rho \cdot OPT_P$.

The value ρ is the **approximation ratio**.

We say that \mathcal{A} achieves an approximation ratio of ρ .

Approximation Algorithm with $\rho = 2$

Consider the following greedy algorithm:

Input: P

1. $C \leftarrow \emptyset$
2. add to C an arbitrary point in P
3. **for** $i = 2$ to k **do**
4. $p \leftarrow$ a point in P with the maximum $dist_C(p)$
5. add p to C
6. return C

The algorithm can be easily implemented in polynomial time.

Later, we will prove that the algorithm is 2-approximate.

Example: $k = 3$




Initially, $C = \{c_1\}$

Example: $k = 3$



After a round, $C = \{c_1, c_2\}$

Example: $k = 3$



After another round, $C = \{c_1, c_2, c_3\}$

Proof of Approximation Guarantee

The approximation guarantee is established by the following theorem.

Theorem 1: The algorithm returns a set C with $\text{pen}(C) \leq 2 \cdot OPT_P$.

Next, we prove this theorem.

Proof: Let $C^* = \{c_1^*, c_2^*, \dots, c_k^*\}$ be an optimal solution, i.e., $\text{pen}(C^*) = OPT_P$.

For each $i \in [1, k]$, define P_i^* as the set of points $p \in P$ satisfying

$$\text{dist}(p, c_i^*) \leq \text{dist}(p, c_j^*)$$

for any $j \neq i$.

(Intuitively, P_i^* is the set of points clustered around the i -th center c_i^* .)

Observation:

For any point $p \in P_i^*$, $\text{dist}(p, c_i^*) = \text{dist}_{C^*}(p) \leq \text{pen}(C^*)$.

Let C_{ours} be the output of our algorithm.

Case 1: C_{ours} has a point in each of $P_1^*, P_2^*, \dots, P_k^*$.

Consider any point $p \in P$. Suppose that $p \in P_i^*$ for some $i \in [1, k]$.

Let c be a point in $C_{ours} \cap P_i^*$. It holds that:

$$\begin{aligned} dist_{C_{ours}}(p) &\leq dist(c, p) && (\text{by def. of distance}) \\ &\leq dist(c, c_i^*) + dist(c_i^*, p) && (\text{by triangle inequality}) \\ &\leq 2 \cdot pen(C^*) && (\text{by def. of penalty}) \end{aligned}$$

Therefore:

$$pen(C_{ours}) = \max_{p \in P} dist_{C_{ours}}(p) \leq 2 \cdot pen(C^*).$$

Case 2: C_{ours} has no point in at least one of P_1^*, \dots, P_k^* . Hence, one of P_1^*, \dots, P_k^* — say P_i^* — must cover at least two points c_a and c_b of C_{ours} . It thus follows that

$$\begin{aligned} dist(c_a, c_b) &\leq dist(c_a, c_i^*) + dist(c_b, c_i^*) \quad (\text{by triangle inequality}) \\ &\leq 2 \cdot pen(C^*). \quad (\text{by definition of penalty}) \end{aligned}$$

Next, we prove:

Claim 1: For any point $p \in P$, $dist_{C_{ours}}(p) \leq dist(c_a, c_b)$.

Note that **Claim 1** implies $pen(C_{ours}) \leq 2 \cdot pen(C^*)$.

This finishes the proof of **Theorem 1** (modulo **Claim 1** which we prove next).

Proof of Claim 1

W.l.o.g., assume that c_b was picked after c_a by our algorithm. Consider the moment right before c_b was picked. At that moment, the set C maintained by our algorithm was a proper subset of C_{ours} .

From the fact that c_b was the next point picked, we know $dist_C(p) \leq dist_C(c_b)$ for every $p \in P$.

Because $c_a \in C$, it holds that $dist_C(c_b) \leq dist(c_a, c_b)$.

The lemma then follows because

$$dist_{C_{ours}}(p) \leq dist_C(p) \leq dist_C(c_b) \leq dist(c_a, c_b).$$

This finishes the proof of **Claim 1**. □